

# ”Music Genre Classification Project”

## Music Genre Classification Project

### Complete Project Code Documentation

**Project:** Music Genre Classification using Machine Learning

**Authors:** Niyat Kahsay & Marwah Faraj

**Date:** Summer 2025

---

This document contains all project code and documentation converted to PDF format using Jupyter nbconvert and pandoc.

**Generated on:** August 08, 2025 at 08:25 PM

---

# 01\_ExploratoryAnalysis\_Feature\_Engineering

August 8, 2025

##Project: Music Genre Classification using Machine Learning

Niyat Kahsay & Marwah Faraj Summer 2025

Description:

This project aims to automatically classify songs into genres based on audio features provided in the Spotify 1.2M Songs Dataset. The workflow includes data exploration, preprocessing, model building, evaluation, and visualization.

Purpose:

Apply supervised learning techniques on real-world audio data Explore audio feature-based genre classification Build a portfolio-ready project demonstrating practical machine learning skills

## 0.1 Import Libraries

```
[ ]: import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import OrdinalEncoder, LabelEncoder, StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn.metrics import mean_squared_error, r2_score, accuracy_score, \
    classification_report, confusion_matrix
from imblearn.over_sampling import SMOTE
from xgboost import XGBClassifier
from sklearn.ensemble import RandomForestClassifier
import numpy as np
import pandas as pd
import plotly.express as px
import plotly.graph_objects as go
from scipy.stats import mode
import warnings
import os
from matplotlib.colors import LinearSegmentedColormap

warnings.filterwarnings("ignore", category=FutureWarning)
```

```
[14]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

## 0.2 Load Data

```
[15]: import pandas as pd
data = pd.read_csv('/content/drive/MyDrive/music-genre-classification/data/
↳spotify_songs.csv')
data.head()
```

```
[15]:
```

	track_id	track_name \
0	6f807x0ima9a1j3VPbc7VN	I Don't Care (with Justin Bieber) - Loud Luxur...
1	Or7CVbZTWZgbTCYdfa2P31	Memories - Dillon Francis Remix
2	1z1Hg7Vb0AhHDiEmnDE79l	All the Time - Don Diablo Remix
3	75FpbthrwQmzHlBJLuGdC7	Call You Mine - Keanu Silva Remix
4	1e8PAfcKUYoKkxPhrHqw4x	Someone You Loved - Future Humans Remix

	track_artist	track_popularity	track_album_id \
0	Ed Sheeran	66	2oCsODGTsR098Gh5ZS12Cx
1	Maroon 5	67	63rPS0264uRjW1X5E6cWv6
2	Zara Larsson	70	1HoSmj2eLcsrR0vE9gThr4
3	The Chainsmokers	60	1nqYs0eflyKKuG0Vchbsk6
4	Lewis Capaldi	69	7m7vv9wlQ4iOLFujIE2zsQ

	track_album_name	track_album_release_date \
0	I Don't Care (with Justin Bieber) [Loud Luxury...	2019-06-14
1	Memories (Dillon Francis Remix)	2019-12-13
2	All the Time (Don Diablo Remix)	2019-07-05
3	Call You Mine - The Remixes	2019-07-19
4	Someone You Loved (Future Humans Remix)	2019-03-05

	playlist_name	playlist_id	playlist_genre	...	key	loudness \
0	Pop Remix	37i9dQZF1DXcZDD7cfEKhW	pop	...	6	-2.634
1	Pop Remix	37i9dQZF1DXcZDD7cfEKhW	pop	...	11	-4.969
2	Pop Remix	37i9dQZF1DXcZDD7cfEKhW	pop	...	1	-3.432
3	Pop Remix	37i9dQZF1DXcZDD7cfEKhW	pop	...	7	-3.778
4	Pop Remix	37i9dQZF1DXcZDD7cfEKhW	pop	...	1	-4.672

	mode	speechiness	acousticness	instrumentalness	liveness	valence \
0	1	0.0583	0.1020	0.000000	0.0653	0.518
1	1	0.0373	0.0724	0.004210	0.3570	0.693
2	0	0.0742	0.0794	0.000023	0.1100	0.613
3	1	0.1020	0.0287	0.000009	0.2040	0.277
4	1	0.0359	0.0803	0.000000	0.0833	0.725

	tempo	duration_ms
0	122.036	194754

```

1    99.972      162600
2   124.008      176616
3   121.956      169093
4   123.976      189052

```

[5 rows x 23 columns]

#Data Exploration

```
[16]: print(data.columns)
```

```

Index(['track_id', 'track_name', 'track_artist', 'track_popularity',
      'track_album_id', 'track_album_name', 'track_album_release_date',
      'playlist_name', 'playlist_id', 'playlist_genre', 'playlist_subgenre',
      'danceability', 'energy', 'key', 'loudness', 'mode', 'speechiness',
      'acousticness', 'instrumentalness', 'liveness', 'valence', 'tempo',
      'duration_ms'],
      dtype='object')

```

```
[17]: data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32833 entries, 0 to 32832
Data columns (total 23 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   track_id                             32833 non-null  object
1   track_name                           32828 non-null  object
2   track_artist                         32828 non-null  object
3   track_popularity                     32833 non-null  int64
4   track_album_id                       32833 non-null  object
5   track_album_name                     32828 non-null  object
6   track_album_release_date             32833 non-null  object
7   playlist_name                        32833 non-null  object
8   playlist_id                          32833 non-null  object
9   playlist_genre                       32833 non-null  object
10  playlist_subgenre                    32833 non-null  object
11  danceability                         32833 non-null  float64
12  energy                              32833 non-null  float64
13  key                                  32833 non-null  int64
14  loudness                             32833 non-null  float64
15  mode                                 32833 non-null  int64
16  speechiness                          32833 non-null  float64
17  acousticness                         32833 non-null  float64
18  instrumentalness                     32833 non-null  float64
19  liveness                             32833 non-null  float64
20  valence                              32833 non-null  float64
21  tempo                                32833 non-null  float64
22  duration_ms                          32833 non-null  int64

```

```
dtypes: float64(9), int64(4), object(10)
memory usage: 5.8+ MB
```

```
[18]: data.describe(include = "all")
```

```
[18]:
```

	track_id	track_name	track_artist	track_popularity	\
count	32833	32828	32828	32833.000000	
unique	28356	23449	10692	NaN	
top	7BKL CZ1jbUBVqRi2FVlTVw	Poison	Martin Garrix	NaN	
freq	10	22	161	NaN	
mean	NaN	NaN	NaN	42.477081	
std	NaN	NaN	NaN	24.984074	
min	NaN	NaN	NaN	0.000000	
25%	NaN	NaN	NaN	24.000000	
50%	NaN	NaN	NaN	45.000000	
75%	NaN	NaN	NaN	62.000000	
max	NaN	NaN	NaN	100.000000	

	track_album_id	track_album_name	track_album_release_date	\
count	32833	32828	32833	
unique	22545	19743	4530	
top	5L1xcowSxwzFUSJzvyMp48	Greatest Hits	2020-01-10	
freq	42	139	270	
mean	NaN	NaN	NaN	
std	NaN	NaN	NaN	
min	NaN	NaN	NaN	
25%	NaN	NaN	NaN	
50%	NaN	NaN	NaN	
75%	NaN	NaN	NaN	
max	NaN	NaN	NaN	

	playlist_name	playlist_id	playlist_genre	...	\
count	32833	32833	32833	...	
unique	449	471	6	...	
top	Indie Poptimism	4JkkvMpVl4lSioqQjeAL0q	edm	...	
freq	308	247	6043	...	
mean	NaN	NaN	NaN	...	
std	NaN	NaN	NaN	...	
min	NaN	NaN	NaN	...	
25%	NaN	NaN	NaN	...	
50%	NaN	NaN	NaN	...	
75%	NaN	NaN	NaN	...	
max	NaN	NaN	NaN	...	

	key	loudness	mode	speechiness	acousticness	\
count	32833.000000	32833.000000	32833.000000	32833.000000	32833.000000	
unique	NaN	NaN	NaN	NaN	NaN	

top	NaN	NaN	NaN	NaN	NaN
freq	NaN	NaN	NaN	NaN	NaN
mean	5.374471	-6.719499	0.565711	0.107068	0.175334
std	3.611657	2.988436	0.495671	0.101314	0.219633
min	0.000000	-46.448000	0.000000	0.000000	0.000000
25%	2.000000	-8.171000	0.000000	0.041000	0.015100
50%	6.000000	-6.166000	1.000000	0.062500	0.080400
75%	9.000000	-4.645000	1.000000	0.132000	0.255000
max	11.000000	1.275000	1.000000	0.918000	0.994000

	instrumentalness	liveness	valence	tempo	\
count	32833.000000	32833.000000	32833.000000	32833.000000	
unique	NaN	NaN	NaN	NaN	
top	NaN	NaN	NaN	NaN	
freq	NaN	NaN	NaN	NaN	
mean	0.084747	0.190176	0.510561	120.881132	
std	0.224230	0.154317	0.233146	26.903624	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.092700	0.331000	99.960000	
50%	0.000016	0.127000	0.512000	121.984000	
75%	0.004830	0.248000	0.693000	133.918000	
max	0.994000	0.996000	0.991000	239.440000	

	duration_ms
count	32833.000000
unique	NaN
top	NaN
freq	NaN
mean	225799.811622
std	59834.006182
min	4000.000000
25%	187819.000000
50%	216000.000000
75%	253585.000000
max	517810.000000

[11 rows x 23 columns]

```
[24]: # Genre Feature Radar Chart (Interactive)
genre_means = data.groupby('playlist_genre')[audio_features].mean().
    ↪reset_index()

fig = go.Figure()
for genre in genre_means['playlist_genre']:
    fig.add_trace(go.Scatterpolar(
        r=genre_means[genre_means['playlist_genre'] == genre][audio_features].
        ↪values[0],
```

```

        theta=audio_features,
        fill='toself',
        name=genre
    ))

fig.update_layout(
    polar=dict(radialaxis=dict(visible=True, range=[0, 1])),
    showlegend=True,
    title='Audio Feature Profiles by Genre',
    height=600
)
fig.show()

```

##Data Preprocessing

```

[25]: # Handle duplicates
data = data.drop_duplicates(subset=['track_id'])

# Fix release dates
def fix_date(x):
    if pd.isnull(x):
        return x
    if isinstance(x, str):
        return f"{x}-01-01" if len(x) < 10 else x
    return x.strftime('%Y-%m-%d')

data['track_album_release_date'] = data['track_album_release_date'].
    ↪apply(fix_date)
data['track_album_release_date'] = pd.
    ↪to_datetime(data['track_album_release_date'], errors='coerce')

# Extract release year
data['release_year'] = data['track_album_release_date'].dt.year

# Encode Target Variable
label_encoder = LabelEncoder()
data['genre_label'] = label_encoder.fit_transform(data['playlist_genre'])
print(" Target encoding complete!")
print(label_encoder.classes_)

# Feature Scaling
audio_features += ['duration_ms', 'release_year']
scaler = StandardScaler()
data[audio_features] = scaler.fit_transform(data[audio_features])

```

```

Target encoding complete!
['edm' 'latin' 'pop' 'r&b' 'rap' 'rock']

```

```
[26]: print(data.isna().sum())
data = data.dropna()
```

```
track_id          0
track_name        4
track_artist      4
track_popularity  0
track_album_id    0
track_album_name  4
track_album_release_date  25
playlist_name     0
playlist_id       0
playlist_genre    0
playlist_subgenre  0
danceability      0
energy            0
key              0
loudness          0
mode             0
speechiness       0
acousticness      0
instrumentalness  0
liveness          0
valence           0
tempo            0
duration_ms       0
release_year      25
genre_label       0
dtype: int64
```

```
[26]:
```

```
[ ]: # Install kaleido for plotly static image export if needed
try:
    import kaleido
except ImportError:
    import subprocess
    import sys
    subprocess.check_call([sys.executable, "-m", "pip", "install", "kaleido"])
    import kaleido

warnings.filterwarnings("ignore", category=FutureWarning)

#
# WHITE BACKGROUND & GOLDEN THEME SETUP
#

# Set up white background and golden color scheme
```



```

plt.style.use('default') # Use default style for white background
golden_palette = ['#8B4513', '#CD853F', '#D4AF37', '#DAA520', '#B8860B',
↳ '#A0522D'] # Adjusted golden palette

print(" Libraries imported successfully!")
print(" White background and golden color scheme activated!")

# Define the base directory for saving plots in Google Drive
PLOT_SAVE_DIR = '/content/drive/MyDrive/music-genre-classification/images/eda'
os.makedirs(PLOT_SAVE_DIR, exist_ok=True)
print(f" Plots will be saved to: {PLOT_SAVE_DIR}")

#
#  UPDATED VISUALIZATION CODE
#

# 1. GENRE DISTRIBUTION VISUALIZATION
def create_genre_distribution_plot(data):
    """Create genre distribution bar chart with white background and golden
↳ aesthetic"""
    plt.figure(figsize=(12, 6), facecolor='white')
    genre_counts = data['playlist_genre'].value_counts()
    ax = sns.barplot(x=genre_counts.index, y=genre_counts.values,
↳ palette=golden_palette)
    ax.set_facecolor('white')
    plt.title('Song Distribution by Genre', fontsize=18, color='black',
↳ fontweight='bold', pad=20)
    plt.xlabel('Genre', fontsize=14, color='black', fontweight='semibold')
    plt.ylabel('Count', fontsize=14, color='black', fontweight='semibold')

    # Style the plot
    ax.tick_params(colors='black', labelsiz=11)
    ax.spines['bottom'].set_color('black')
    ax.spines['left'].set_color('black')
    ax.spines['top'].set_visible(False)
    ax.spines['right'].set_visible(False)
    ax.grid(True, alpha=0.2, color='gray')

    # Add percentage labels
    total = len(data)
    for p in ax.patches:
        percentage = f'{100 * p.get_height()/total:.1f}%'
        ax.annotate(percentage, (p.get_x() + p.get_width()/2., p.get_height()),
            ha='center', va='center', xytext=(0, 10),
            textcoords='offset points', fontsize=11, color='black',
↳ fontweight='bold')

```

```

# Save the plot
plt.savefig(os.path.join(PLOT_SAVE_DIR, 'genre_distribution.png'),
            facecolor='white', edgecolor='none', dpi=300,
↳bbox_inches='tight')
plt.show()
print(f" Plot saved to: {os.path.join(PLOT_SAVE_DIR, 'genre_distribution.
↳png')}")

# 2. CORRELATION MATRIX VISUALIZATION
def create_correlation_matrix(data):
    """Create correlation matrix heatmap with white background and golden
↳theme"""
    plt.figure(figsize=(14, 10), facecolor='white')
    audio_features = ['danceability', 'energy', 'loudness',
                      'acousticness', 'valence', 'tempo',
                      'speechiness', 'instrumentalness']
    corr_matrix = data[audio_features].corr()
    mask = np.triu(np.ones_like(corr_matrix, dtype=bool))

    # Create custom colormap with golden tones
    colors = ['white', '#F5DEB3', '#D4AF37', '#CD853F', '#8B4513', '#2C1810']
    golden_cmap = LinearSegmentedColormap.from_list('golden', colors, N=256)

    ax = plt.gca()
    ax.set_facecolor('white')
    sns.heatmap(corr_matrix, mask=mask, annot=True, fmt=".2f",
                cmap=golden_cmap, linewidths=0.5, linecolor='gray',
                annot_kws={'color': 'black', 'fontweight': 'bold'},
                cbar_kws={'shrink': 0.8})
    plt.title("Audio Feature Correlation Matrix", fontsize=18, color='black',
↳fontweight='bold', pad=20)

# Style axes
ax.tick_params(colors='black', labelsize=11)
plt.xticks(rotation=45, ha='right', color='black')
plt.yticks(rotation=0, color='black')

# Save the plot
plt.savefig(os.path.join(PLOT_SAVE_DIR, 'correlation_matrix.png'),
            facecolor='white', edgecolor='none', dpi=300,
↳bbox_inches='tight')
plt.show()
print(f" Plot saved to: {os.path.join(PLOT_SAVE_DIR, 'correlation_matrix.
↳png')}")

```

```

# 3. FEATURE DISTRIBUTION BOXPLOTS
def create_feature_boxplots(data):
    """Create feature distribution boxplots with white background and golden
    ↪theme"""
    audio_features = ['danceability', 'energy', 'loudness',
                      'acousticness', 'valence', 'tempo',
                      'speechiness', 'instrumentalness']

    plt.figure(figsize=(15, 10), facecolor='white')
    for i, feature in enumerate(audio_features[:6], 1):
        plt.subplot(2, 3, i)
        ax = plt.gca()
        ax.set_facecolor('white')

        sns.boxplot(x='playlist_genre', y=feature, data=data,
        ↪palette=golden_palette)
        plt.title(f'{feature.capitalize()} by Genre', fontsize=14,
        ↪color='black', fontweight='bold', pad=10)
        plt.xticks(rotation=45, color='black', fontsize=10)
        plt.yticks(color='black', fontsize=10)
        plt.xlabel('Genre', color='black', fontsize=11, fontweight='semibold')
        plt.ylabel(feature.capitalize(), color='black', fontsize=11,
        ↪fontweight='semibold')

        # Style the subplot
        ax.tick_params(colors='black')
        for spine in ax.spines.values():
            spine.set_color('black')
            spine.set_linewidth(0.8)
        ax.grid(True, alpha=0.2, color='gray')

    plt.tight_layout()
    # Save the plot
    plt.savefig(os.path.join(PLOT_SAVE_DIR, 'feature_distributions.png'),
                facecolor='white', edgecolor='none', dpi=300,
    ↪bbox_inches='tight')
    plt.show()
    print(f" Plot saved to: {os.path.join(PLOT_SAVE_DIR,
    ↪'feature_distributions.png')}")

# 4. INTERACTIVE RADAR CHART
def create_radar_chart(data):
    """Create interactive radar chart with white background and golden theme"""
    audio_features = ['danceability', 'energy', 'loudness',
                      'acousticness', 'valence', 'tempo',
                      'speechiness', 'instrumentalness']

```

```

genre_means = data.groupby('playlist_genre')[audio_features].mean().
↳reset_index()

# Define golden color palette for radar chart
radar_colors = ['#8B4513', '#CD853F', '#D4AF37', '#DAA520', '#B8860B',
↳'#A0522D'] # Adjusted golden palette

fig = go.Figure()
for i, genre in enumerate(genre_means['playlist_genre']):
    fig.add_trace(go.Scatterpolar(
        r=genre_means[genre_means['playlist_genre'] ==
↳genre][audio_features].values[0],
        theta=audio_features,
        fill='toself',
        name=genre,
        line=dict(color=radar_colors[i % len(radar_colors)], width=3),
        fillcolor=radar_colors[i % len(radar_colors)],
        opacity=0.6
    ))

fig.update_layout(
    polar=dict(
        radialaxis=dict(
            visible=True,
            range=[0, 1],
            tickfont=dict(color='black', size=12),
            gridcolor='gray',
            linecolor='gray'
        ),
        angularaxis=dict(
            tickfont=dict(color='black', size=12),
            gridcolor='gray',
            linecolor='gray'
        ),
        bgcolor='white'
    ),
    showlegend=True,
    title=dict(
        text='Audio Feature Profiles by Genre',
        font=dict(color='black', size=18, family='Arial Black'),
        x=0.5
    ),
    height=600,
    paper_bgcolor='white',
    plot_bgcolor='white',
    font=dict(color='black'),

```

```

        legend=dict(
            font=dict(color='black', size=12),
            bgcolor='rgba(255, 255, 255, 0.8)',
            bordercolor='gray',
            borderwidth=1
        )
    )

    # Save the interactive plot as HTML
    fig.write_html(os.path.join(PLOT_SAVE_DIR, 'radar_chart.html'))
    # fig.write_image(os.path.join(PLOT_SAVE_DIR, 'radar_chart.png'),
    ↪width=800, height=600) # Still commenting this out due to kaleido issues
    fig.show()
    print(f" Interactive plot saved to: {os.path.join(PLOT_SAVE_DIR,
    ↪'radar_chart.html')}")
    # print(f" Static plot saved to: {os.path.join(PLOT_SAVE_DIR, 'radar_chart.
    ↪png')}")

#
#  USAGE INSTRUCTIONS
#

"""
USAGE INSTRUCTIONS:

1. Replace your existing plotting cells in the notebook with the function calls:

    # Instead of the old genre distribution code, use:
    create_genre_distribution_plot(data)

    # Instead of the old correlation matrix code, use:
    create_correlation_matrix(data)

    # Instead of the old boxplot code, use:
    create_feature_boxplots(data)

    # Instead of the old radar chart code, use:
    create_radar_chart(data)

2. All plots will automatically be saved to the directory specified by
    ↪PLOT_SAVE_DIR

3. The color scheme matches your presentation slide with:
    - Dark brown background (#2C1810)
    - Golden palette for data visualization
    - Elegant typography and styling

```

```

4. Make sure your data loading path is correct for your environment:
- Colab: '/content/spotify_songs.csv'
- Local: './data/spotify_songs.csv'
- Current: '/workspace/data/spotify_songs.csv'
"""

print(" Updated Feature Engineering Code Ready!")
print(" Copy the functions above into your notebook cells")
print(" Your visualizations will match the slide aesthetic perfectly!")


```

Libraries imported successfully!  
 White background and golden color scheme activated!  
 Plots will be saved to: /content/drive/MyDrive/music-genre-  
 classification/images/eda  
 Updated Feature Engineering Code Ready!  
 Copy the functions above into your notebook cells  
 Your visualizations will match the slide aesthetic perfectly!


```

[38]: # Call the plotting functions to generate and save plots
create_genre_distribution_plot(data)
create_correlation_matrix(data)
create_feature_boxplots(data)
create_radar_chart(data)

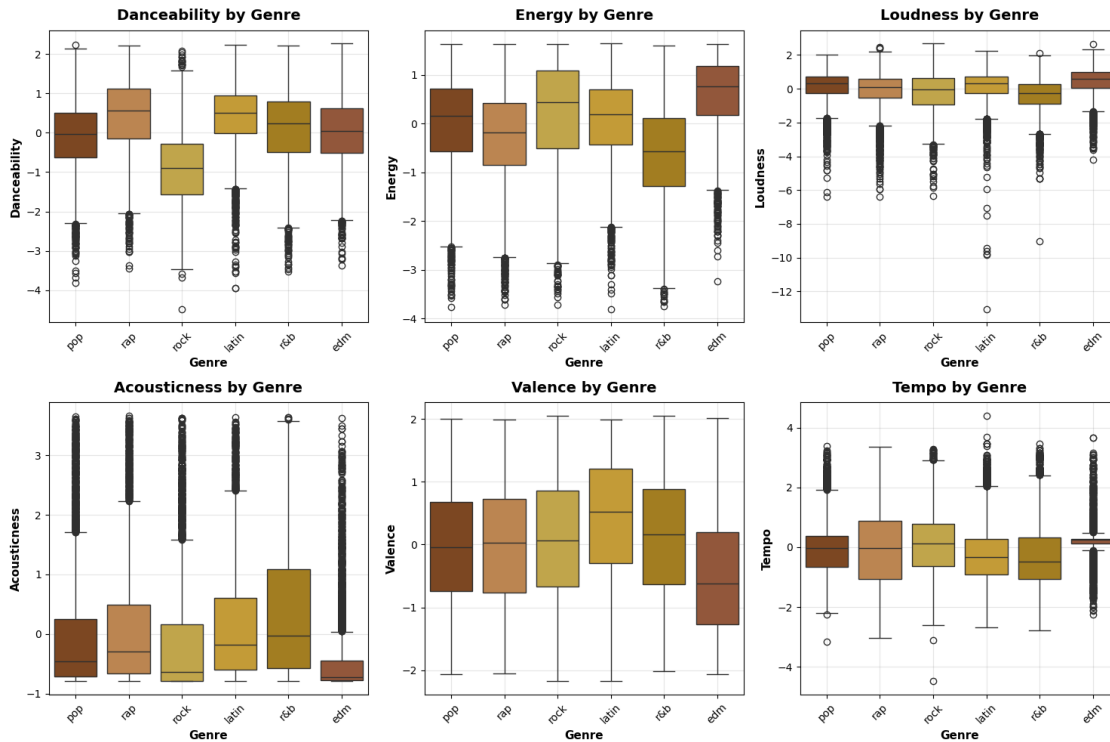
```



Plot saved to: /content/drive/MyDrive/music-genre-  
 classification/images/eda/genre\_distribution.png



Plot saved to: /content/drive/MyDrive/music-genre-classification/images/eda/correlation\_matrix.png



Plot saved to: /content/drive/MyDrive/music-genre-classification/images/eda/feature\_distributions.png

Interactive plot saved to: /content/drive/MyDrive/music-genre-classification/images/eda/radar\_chart.html

```
[31]: %pip install -U kaleido
```

Requirement already satisfied: kaleido in /usr/local/lib/python3.11/dist-packages (1.0.0)

Requirement already satisfied: choreographer>=1.0.5 in /usr/local/lib/python3.11/dist-packages (from kaleido) (1.0.9)

Requirement already satisfied: logistro>=1.0.8 in /usr/local/lib/python3.11/dist-packages (from kaleido) (1.1.0)

Requirement already satisfied: orjson>=3.10.15 in /usr/local/lib/python3.11/dist-packages (from kaleido) (3.11.1)

Requirement already satisfied: packaging in /usr/local/lib/python3.11/dist-packages (from kaleido) (25.0)

Requirement already satisfied: simplejson>=3.19.3 in /usr/local/lib/python3.11/dist-packages (from choreographer>=1.0.5->kaleido) (3.20.1)



# 02\_Modeling

August 8, 2025

##Project: Music Genre Classification using Machine Learning

Niyat Kahsay & Marwah Faraj Summer 2025

Description:

This project aims to automatically classify songs into genres based on audio features provided in the Spotify 1.2M Songs Dataset. The workflow includes data exploration, preprocessing, model building, evaluation, and visualization.

Purpose:

Apply supervised learning techniques on real-world audio data Explore audio feature-based genre classification Build a portfolio-ready project demonstrating practical machine learning skills

#Data Preparation

## 0.1 Import Libraries

```
[1]: import seaborn as sns
import matplotlib.pyplot as plt
import plotly.graph_objects as go
from sklearn.preprocessing import OrdinalEncoder
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report
from sklearn.metrics import confusion_matrix
from imblearn.over_sampling import SMOTE
from xgboost import XGBClassifier
from sklearn.metrics import (accuracy_score, confusion_matrix,
                             classification_report, f1_score)
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from xgboost import XGBClassifier
from sklearn.model_selection import cross_val_score, GridSearchCV, train_test_split
from sklearn.svm import SVC
```

```

from sklearn.decomposition import PCA
import numpy as np
import time
import pandas as pd
import plotly.express as px
from scipy.stats import mode
import warnings

warnings.filterwarnings("ignore", category=FutureWarning)

```

## 0.2 Load Data

```

[2]: from google.colab import drive
drive.mount('/content/drive')

```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

```

[2]: import pandas as pd
# data = pd.read_csv('/content/drive/MyDrive/music-genre-classification/data/
↳spotify_songs.csv')
data= pd.read_csv('/Users/marwahfaraj/Desktop/ms_degree_application_and_doc/
↳final_projects/504_final_project/music-genre-classification/data/
↳spotify_songs.csv')
data.head()

```

```

[2]:

```

	track_id	track_name \
0	6f807x0ima9a1j3VPbc7VN	I Don't Care (with Justin Bieber) - Loud Luxur...
1	0r7CVbZTWZgbTCYdfa2P31	Memories - Dillon Francis Remix
2	1z1Hg7Vb0AhHDE791	All the Time - Don Diablo Remix
3	75FpbthrwQmzH1BJLuGdC7	Call You Mine - Keanu Silva Remix
4	1e8PAfckUYoKkxPhrHqw4x	Someone You Loved - Future Humans Remix

	track_artist	track_popularity	track_album_id \
0	Ed Sheeran	66	2oCsODGTsR098Gh5ZS12Cx
1	Maroon 5	67	63rPS0264uRjW1X5E6cWv6
2	Zara Larsson	70	1HoSmj2eLcsrR0vE9gThr4
3	The Chainsmokers	60	1nqYs0eflyKKuG0Vchbsk6
4	Lewis Capaldi	69	7m7vv9wlQ4i0LFuJiE2zsQ

	track_album_name	track_album_release_date \
0	I Don't Care (with Justin Bieber) [Loud Luxury...	2019-06-14
1	Memories (Dillon Francis Remix)	2019-12-13
2	All the Time (Don Diablo Remix)	2019-07-05
3	Call You Mine - The Remixes	2019-07-19
4	Someone You Loved (Future Humans Remix)	2019-03-05

	playlist_name	playlist_id	playlist_genre	...	key	loudness	\
0	Pop Remix	37i9dQZF1DXcZDD7cfEKhW	pop	...	6	-2.634	
1	Pop Remix	37i9dQZF1DXcZDD7cfEKhW	pop	...	11	-4.969	
2	Pop Remix	37i9dQZF1DXcZDD7cfEKhW	pop	...	1	-3.432	
3	Pop Remix	37i9dQZF1DXcZDD7cfEKhW	pop	...	7	-3.778	
4	Pop Remix	37i9dQZF1DXcZDD7cfEKhW	pop	...	1	-4.672	

	mode	speechiness	acousticness	instrumentalness	liveness	valence	\
0	1	0.0583	0.1020	0.000000	0.0653	0.518	
1	1	0.0373	0.0724	0.004210	0.3570	0.693	
2	0	0.0742	0.0794	0.000023	0.1100	0.613	
3	1	0.1020	0.0287	0.000009	0.2040	0.277	
4	1	0.0359	0.0803	0.000000	0.0833	0.725	

	tempo	duration_ms
0	122.036	194754
1	99.972	162600
2	124.008	176616
3	121.956	169093
4	123.976	189052

[5 rows x 23 columns]

#Data Exploration

```
[3]: print(data.columns)
```

```
Index(['track_id', 'track_name', 'track_artist', 'track_popularity',
      'track_album_id', 'track_album_name', 'track_album_release_date',
      'playlist_name', 'playlist_id', 'playlist_genre', 'playlist_subgenre',
      'danceability', 'energy', 'key', 'loudness', 'mode', 'speechiness',
      'acousticness', 'instrumentalness', 'liveness', 'valence', 'tempo',
      'duration_ms'],
      dtype='object')
```

```
[4]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 32833 entries, 0 to 32832
```

```
Data columns (total 23 columns):
```

#	Column	Non-Null Count	Dtype
0	track_id	32833 non-null	object
1	track_name	32828 non-null	object
2	track_artist	32828 non-null	object
3	track_popularity	32833 non-null	int64
4	track_album_id	32833 non-null	object
5	track_album_name	32828 non-null	object
6	track_album_release_date	32833 non-null	object

```

7  playlist_name          32833 non-null object
8  playlist_id            32833 non-null object
9  playlist_genre         32833 non-null object
10 playlist_subgenre      32833 non-null object
11 danceability           32833 non-null float64
12 energy                 32833 non-null float64
13 key                    32833 non-null int64
14 loudness               32833 non-null float64
15 mode                   32833 non-null int64
16 speechiness            32833 non-null float64
17 acousticness           32833 non-null float64
18 instrumentalness       32833 non-null float64
19 liveness               32833 non-null float64
20 valence                 32833 non-null float64
21 tempo                  32833 non-null float64
22 duration_ms            32833 non-null int64

```

dtypes: float64(9), int64(4), object(10)

memory usage: 5.8+ MB

```
[5]: data.describe(include = "all")
```

```

[5]:
count          track_id track_name  track_artist  track_popularity \
unique          28356    23449         10692             NaN
top      7BKLCZ1jbUBVqRi2FVlTVw    Poison  Martin Garrix             NaN
freq           10         22          161             NaN
mean           NaN         NaN         NaN      42.477081
std            NaN         NaN         NaN      24.984074
min            NaN         NaN         NaN       0.000000
25%            NaN         NaN         NaN      24.000000
50%            NaN         NaN         NaN      45.000000
75%            NaN         NaN         NaN      62.000000
max            NaN         NaN         NaN     100.000000

```

```

count      track_album_id track_album_name track_album_release_date \
unique      22545         19743             4530
top      5L1xcowSxwzFUSJzvyMp48    Greatest Hits      2020-01-10
freq           42          139             270
mean           NaN         NaN         NaN
std            NaN         NaN         NaN
min            NaN         NaN         NaN
25%            NaN         NaN         NaN
50%            NaN         NaN         NaN
75%            NaN         NaN         NaN
max            NaN         NaN         NaN

```

	playlist_name	playlist_id	playlist_genre	...	\
count	32833	32833	32833	...	
unique	449	471	6	...	
top	Indie Poptimism	4JkkvMpVl4lSioqQjeAL0q	edm	...	
freq	308	247	6043	...	
mean	NaN	NaN	NaN	...	
std	NaN	NaN	NaN	...	
min	NaN	NaN	NaN	...	
25%	NaN	NaN	NaN	...	
50%	NaN	NaN	NaN	...	
75%	NaN	NaN	NaN	...	
max	NaN	NaN	NaN	...	

	key	loudness	mode	speechiness	acousticness	\
count	32833.000000	32833.000000	32833.000000	32833.000000	32833.000000	
unique	NaN	NaN	NaN	NaN	NaN	
top	NaN	NaN	NaN	NaN	NaN	
freq	NaN	NaN	NaN	NaN	NaN	
mean	5.374471	-6.719499	0.565711	0.107068	0.175334	
std	3.611657	2.988436	0.495671	0.101314	0.219633	
min	0.000000	-46.448000	0.000000	0.000000	0.000000	
25%	2.000000	-8.171000	0.000000	0.041000	0.015100	
50%	6.000000	-6.166000	1.000000	0.062500	0.080400	
75%	9.000000	-4.645000	1.000000	0.132000	0.255000	
max	11.000000	1.275000	1.000000	0.918000	0.994000	

	instrumentalness	liveness	valence	tempo	\
count	32833.000000	32833.000000	32833.000000	32833.000000	
unique	NaN	NaN	NaN	NaN	
top	NaN	NaN	NaN	NaN	
freq	NaN	NaN	NaN	NaN	
mean	0.084747	0.190176	0.510561	120.881132	
std	0.224230	0.154317	0.233146	26.903624	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.092700	0.331000	99.960000	
50%	0.000016	0.127000	0.512000	121.984000	
75%	0.004830	0.248000	0.693000	133.918000	
max	0.994000	0.996000	0.991000	239.440000	

	duration_ms
count	32833.000000
unique	NaN
top	NaN
freq	NaN
mean	225799.811622
std	59834.006182
min	4000.000000

```

25%      187819.000000
50%      216000.000000
75%      253585.000000
max       517810.000000

```

```
[11 rows x 23 columns]
```

```

[6]: # Feature engineering
data['track_album_release_year'] = pd.to_datetime(
    data['track_album_release_date'], errors='coerce').dt.year
data['track_album_release_year'].fillna(data['track_album_release_year'].
    ↳median(), inplace=True)
data['duration_s'] = data['duration_ms'] / 1000 # Convert to seconds
data.drop('duration_ms', axis=1, inplace=True)

```

```

[7]: # Encode target variable
genre_encoder = LabelEncoder()
data['genre_encoded'] = genre_encoder.fit_transform(data['playlist_genre'])

```

```

[8]: # Select relevant features
features = ['danceability', 'energy', 'key', 'loudness', 'mode',
            'speechiness', 'acousticness', 'instrumentalness',
            'liveness', 'valence', 'tempo', 'duration_s',
            'track_popularity', 'track_album_release_year']

X = data[features]
y = data['genre_encoded']

```

```

[9]: # Split data
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y)

```

```

[10]: # Scale features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

```

```

[11]: # Save the fitted scaler for deployment
import joblib
import os
scaler_save_path = '/Users/marwahfaraj/Desktop/ms_degree_application_and_doc/
    ↳final_projects/504_final_project/music-genre-classification/app/model/
    ↳standard_scaler.pkl'
os.makedirs(os.path.dirname(scaler_save_path), exist_ok=True)
joblib.dump(scaler, scaler_save_path)
print(f"StandardScaler saved to: {scaler_save_path}")

```

```
StandardScaler saved to: /Users/marwahfaraj/Desktop/ms_degree_application_and_do
```

c/final\_projects/504\_final\_project/music-genre-classification/app/model/standard\_scaler.pkl

##Model Development

```
[12]: models = {  
    'Random Forest': RandomForestClassifier(random_state=42),  
    'Gradient Boosting': GradientBoostingClassifier(random_state=42),  
    'XGBoost': XGBClassifier(random_state=42, use_label_encoder=False,  
    ↪eval_metric='mlogloss'),  
    'SVM': SVC(random_state=42, probability=True)  
}
```

```
[13]: # Cross-validation evaluation  
results = {}  
for name, model in models.items():  
    start_time = time.time()  
    cv_scores = cross_val_score(model, X_train_scaled, y_train, cv=5,  
    ↪scoring='accuracy')  
    results[name] = {  
        'cv_accuracy': np.mean(cv_scores),  
        'cv_time': time.time() - start_time  
    }  
    print(f"{name} - Avg CV Accuracy: {np.mean(cv_scores):.4f} - Time:␣  
    ↪{results[name]['cv_time']:.2f}s")
```

Random Forest - Avg CV Accuracy: 0.5747 - Time: 44.56s

Gradient Boosting - Avg CV Accuracy: 0.5715 - Time: 284.02s

/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:

[02:03:59] WARNING: /workspace/src/learner.cc:738:

Parameters: { "use\_label\_encoder" } are not used.

bst.update(dtrain, iteration=i, fobj=obj)

/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:

[02:04:02] WARNING: /workspace/src/learner.cc:738:

Parameters: { "use\_label\_encoder" } are not used.

bst.update(dtrain, iteration=i, fobj=obj)

/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:

[02:04:05] WARNING: /workspace/src/learner.cc:738:

Parameters: { "use\_label\_encoder" } are not used.

bst.update(dtrain, iteration=i, fobj=obj)

/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:

[02:04:11] WARNING: /workspace/src/learner.cc:738:

Parameters: { "use\_label\_encoder" } are not used.

bst.update(dtrain, iteration=i, fobj=obj)

```
/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:
[02:04:14] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
```

```
XGBoost - Avg CV Accuracy: 0.5876 - Time: 18.55s
```

```
SVM - Avg CV Accuracy: 0.5548 - Time: 812.40s
```

```
[14]: # Select best model based on CV
best_model_name = max(results, key=lambda x: results[x]['cv_accuracy'])
print(f"\nBest model from CV: {best_model_name}")
```

Best model from CV: XGBoost

```
[15]: # Hyperparameter tuning for best model
# Best model: XGBoost
param_grid = {
    'n_estimators': [100, 200],
    'learning_rate': [0.05, 0.1],
    'max_depth': [3, 5],
    'subsample': [0.8, 1.0]
}

grid_search = GridSearchCV(
    models['XGBoost'],
    param_grid,
    cv=3,
    scoring='accuracy',
    n_jobs=-1,
    verbose=1
)

grid_search.fit(X_train_scaled, y_train)

# Get best model
best_model = grid_search.best_estimator_
print(f"Best parameters: {grid_search.best_params_}")
print(f"Best CV accuracy: {grid_search.best_score_:.4f}")
```

Fitting 3 folds for each of 16 candidates, totalling 48 fits

```
/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:
[02:20:01] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
```

```
Best parameters: {'learning_rate': 0.1, 'max_depth': 5, 'n_estimators': 200,
```



```
'subsample': 0.8}
Best CV accuracy: 0.5883
```

```
##Model Evaluation
```

```
[16]: # Train final model
best_model.fit(X_train_scaled, y_train)

# Predictions
y_pred = best_model.predict(X_test_scaled)
y_proba = best_model.predict_proba(X_test_scaled)

# Evaluation metrics
accuracy = accuracy_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred, average='weighted')
class_report = classification_report(y_test, y_pred, target_names=genre_encoder.
    ↪classes_)

print("\n" + "="*50)
print(f"Final Model: {best_model_name}")
print(f"Test Accuracy: {accuracy:.4f}")
print(f"Weighted F1 Score: {f1:.4f}")
print("\nClassification Report:")
print(class_report)
```

```
/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:
[02:20:09] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
```

```
=====
```

```
Final Model: XGBoost
Test Accuracy: 0.5959
Weighted F1 Score: 0.5940
```

```
Classification Report:
```

	precision	recall	f1-score	support
edm	0.69	0.70	0.70	1209
latin	0.53	0.48	0.51	1031
pop	0.43	0.45	0.44	1102
r&b	0.55	0.48	0.51	1086
rap	0.61	0.67	0.64	1149
rock	0.76	0.79	0.77	990
accuracy			0.60	6567
macro avg	0.59	0.60	0.59	6567

weighted avg	0.59	0.60	0.59	6567
--------------	------	------	------	------

Low performance in 'latin' and 'pop'. Will do, resampling

```
[17]: from imblearn.over_sampling import SMOTE
      # Apply SMOTE to training data
      smote = SMOTE(random_state=42)
      X_train_resampled, y_train_resampled = smote.fit_resample(X_train_scaled,
      ↪ y_train)
```

```
[18]: grid_search_smote = GridSearchCV(
      models['XGBoost'], # Same XGBoost model
      param_grid,
      cv=3,
      scoring='accuracy',
      n_jobs=-1,
      verbose=1
    )

    # Fit the model on the SMOTE-balanced data
    grid_search_smote.fit(X_train_resampled, y_train_resampled)

    # Predict on the original test set (not resampled!)
    y_pred_smote = grid_search_smote.predict(X_test_scaled)
```

Fitting 3 folds for each of 16 candidates, totalling 48 fits

```
/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:
[02:22:42] WARNING: /workspace/src/learner.cc:738:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
```

```
[19]: from sklearn.metrics import accuracy_score, f1_score, classification_report

      print("\n" + "="*50)
      print("SMOTE Model Performance:")
      print("Test Accuracy:", accuracy_score(y_test, y_pred_smote))
      print("Weighted F1 Score:", f1_score(y_test, y_pred_smote, average='weighted'))
      print("\nClassification Report:\n", classification_report(y_test, y_pred_smote))
```

```
=====
SMOTE Model Performance:
Test Accuracy: 0.596771737475255
Weighted F1 Score: 0.5949522484302238
```

```
Classification Report:
      precision    recall  f1-score   support
```

0	0.71	0.69	0.70	1209
1	0.52	0.51	0.51	1031
2	0.43	0.44	0.44	1102
3	0.54	0.47	0.50	1086
4	0.62	0.67	0.64	1149
5	0.74	0.80	0.77	990
accuracy			0.60	6567
macro avg	0.59	0.60	0.59	6567
weighted avg	0.59	0.60	0.59	6567


##But some results went down, so we are keeping the original model

```
[49]: import os
import matplotlib.pyplot as plt # Import pyplot
import seaborn as sns # Import seaborn if needed
import numpy as np # Import numpy if needed
from sklearn.metrics import confusion_matrix # Import confusion_matrix if needed

# Ensure the directory exists
if not os.path.exists('images/models'):
    os.makedirs('images/models')

# Set up white background and golden color scheme
plt.style.use('default') # Use default style for white background
golden_palette = ['#8B4513', '#CD853F', '#D4AF37', '#DAA520', '#B8860B',
                  ↪ '#A0522D'] # Adjusted golden palette

# Confusion matrix
plt.figure(figsize=(12, 10))
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap=golden_palette,
            xticklabels=genre_encoder.classes_,
            yticklabels=genre_encoder.classes_)
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.xticks(rotation=45)
plt.yticks(rotation=0)
plt.tight_layout()
plt.savefig('images/models/confusion_matrix.png')
plt.show()
```



```
[21]: import os
import matplotlib.pyplot as plt # Import pyplot
import numpy as np # Import numpy if needed

# Ensure the directory exists
if not os.path.exists('images/models'):
    os.makedirs('images/models')


# Set up white background and golden color scheme
plt.style.use('default') # Use default style for white background
golden_palette = ['#8B4513', '#CD853F', '#D4AF37', '#DAA520', '#B8860B',
    ↪ '#A0522D'] # Adjusted golden palette

# Feature importance
if hasattr(best_model, 'feature_importances_'):
```

```

plt.figure(figsize=(12, 8))
importances = best_model.feature_importances_
indices = np.argsort(importances)[::-1]
plt.title('Feature Importances')
# Use all colors in the palette, repeating if necessary
colors = np.tile(golden_palette, len(indices) // len(golden_palette) + 1)[:
↪len(indices)]
plt.barh(range(len(indices)), importances[indices], align='center',
↪color=colors) # Changed color to use the palette
plt.yticks(range(len(indices)), [features[i] for i in indices])
plt.gca().invert_yaxis()
plt.tight_layout()
plt.savefig('images/models/feature_importances.png')
plt.show()

```



```

[22]: import os
import matplotlib.pyplot as plt # Import pyplot
import numpy as np # Import numpy if needed
from sklearn.decomposition import PCA # Import PCA if needed

# Ensure the directory exists
if not os.path.exists('images/models'):
    os.makedirs('images/models')

```


```

# Set up white background and golden color scheme
plt.style.use('default') # Use default style for white background
golden_palette = ['#8B4513', '#CD853F', '#D4AF37', '#DAA520', '#B8860B', '#A0522D'] # Adjusted golden palette

# PCA for dimensionality reduction (for visualization)
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_train_scaled)

plt.figure(figsize=(10, 6))
for i, genre in enumerate(np.unique(y_train)):
    plt.scatter(X_pca[y_train == genre, 0],
                X_pca[y_train == genre, 1],
                label=genre_encoder.inverse_transform([genre])[0],
                alpha=0.6,
                color=golden_palette[i % len(golden_palette)])
plt.title('PCA Visualization of Audio Features')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.legend(title='Genre')
plt.grid(alpha=0.3)
plt.tight_layout()
plt.savefig('images/models/pca_visualization.png')
plt.show()

```



## 0.3 Define the deep learning model

### 0.3.1 Subtask:

Generate code to define a suitable deep learning model architecture for music genre classification, considering the nature of your features. This might involve layers like Dense, Dropout, etc.

**Reasoning:** Define a deep learning model architecture for music genre classification using Keras Sequential model, dense layers with ReLU activation, dropout layers, and a final dense layer with softmax activation.

```
[23]: from tensorflow.keras.models import Sequential
      from tensorflow.keras.layers import Dense, Dropout
      import numpy as np

      # Determine the number of unique genres
      num_genres = len(np.unique(y_train))

      # Define the deep learning model
      dl_model = Sequential()

      # Input layer and first hidden layer with dropout
      dl_model.add(Dense(128, activation='relu', input_shape=(X_train_scaled.
        ↳shape[1],)))
      dl_model.add(Dropout(0.3)) # Added dropout

      # Second hidden layer with dropout
      dl_model.add(Dense(64, activation='relu'))
      dl_model.add(Dropout(0.3)) # Added dropout

      # Output layer
      dl_model.add(Dense(num_genres, activation='softmax'))

      # Compile the model
      dl_model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
        ↳metrics=['accuracy']) # Added compile

      dl_model.summary()
```

```
/usr/local/lib/python3.11/dist-packages/keras/src/layers/core/dense.py:93:
UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When
using Sequential models, prefer using an `Input(shape)` object as the first
layer in the model instead.
```

```
    super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 128)	1,920
dropout (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 64)	8,256
dropout_1 (Dropout)	(None, 64)	0
dense_2 (Dense)	(None, 6)	390

Total params: 10,566 (41.27 KB)

Trainable params: 10,566 (41.27 KB)

Non-trainable params: 0 (0.00 B)

## 0.4 Compile the model

### 0.4.1 Subtask:

Add code to compile the deep learning model with an appropriate optimizer, loss function, and metrics.

**Reasoning:** Compile the deep learning model with the specified optimizer, loss function, and metrics.

```
[24]: dl_model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',  
    ↪ metrics=['accuracy'])
```

## 0.5 Train the model

### 0.5.1 Subtask:

Generate code to train the deep learning model on your scaled training data, including considerations for epochs, batch size, and validation data.

**Reasoning:** Train the deep learning model on the scaled training data with specified epochs, batch size, and validation split.

```
[25]: from tensorflow.keras.callbacks import EarlyStopping # Import EarlyStopping  
  
    # Instantiate EarlyStopping callback  
    early_stopping = EarlyStopping(monitor='val_loss', patience=15,  
    ↪ restore_best_weights=True) # Increased patience
```



```

history = dl_model.fit(X_train_scaled, y_train,
                       epochs=50,
                       batch_size=64,
                       validation_split=0.2,
                       callbacks=[early_stopping]) # Added EarlyStopping
↪ callback

```

```

Epoch 1/50
329/329          3s 4ms/step -
accuracy: 0.3627 - loss: 1.5625 - val_accuracy: 0.5010 - val_loss: 1.2921
Epoch 2/50
329/329          2s 3ms/step -
accuracy: 0.4786 - loss: 1.3487 - val_accuracy: 0.5190 - val_loss: 1.2530
Epoch 3/50
329/329          1s 3ms/step -
accuracy: 0.4942 - loss: 1.3086 - val_accuracy: 0.5259 - val_loss: 1.2343
Epoch 4/50
329/329          1s 3ms/step -
accuracy: 0.5067 - loss: 1.2884 - val_accuracy: 0.5306 - val_loss: 1.2171
Epoch 5/50
329/329          1s 3ms/step -
accuracy: 0.5164 - loss: 1.2632 - val_accuracy: 0.5344 - val_loss: 1.2060
Epoch 6/50
329/329          1s 3ms/step -
accuracy: 0.5224 - loss: 1.2604 - val_accuracy: 0.5362 - val_loss: 1.2024
Epoch 7/50
329/329          1s 3ms/step -
accuracy: 0.5287 - loss: 1.2374 - val_accuracy: 0.5430 - val_loss: 1.1877
Epoch 8/50
329/329          1s 3ms/step -
accuracy: 0.5331 - loss: 1.2301 - val_accuracy: 0.5476 - val_loss: 1.1807
Epoch 9/50
329/329          1s 3ms/step -
accuracy: 0.5326 - loss: 1.2252 - val_accuracy: 0.5497 - val_loss: 1.1796
Epoch 10/50
329/329          2s 4ms/step -
accuracy: 0.5449 - loss: 1.2090 - val_accuracy: 0.5495 - val_loss: 1.1725
Epoch 11/50
329/329          1s 4ms/step -
accuracy: 0.5415 - loss: 1.2120 - val_accuracy: 0.5525 - val_loss: 1.1675
Epoch 12/50
329/329          2s 3ms/step -
accuracy: 0.5428 - loss: 1.2024 - val_accuracy: 0.5539 - val_loss: 1.1651
Epoch 13/50
329/329          1s 3ms/step -
accuracy: 0.5402 - loss: 1.2043 - val_accuracy: 0.5586 - val_loss: 1.1566
Epoch 14/50

```

329/329                    1s 3ms/step -  
 accuracy: 0.5409 - loss: 1.2005 - val\_accuracy: 0.5594 - val\_loss: 1.1583  
 Epoch 15/50  
 329/329                    1s 3ms/step -  
 accuracy: 0.5453 - loss: 1.1941 - val\_accuracy: 0.5575 - val\_loss: 1.1530  
 Epoch 16/50  
 329/329                    1s 3ms/step -  
 accuracy: 0.5465 - loss: 1.1893 - val\_accuracy: 0.5638 - val\_loss: 1.1497  
 Epoch 17/50  
 329/329                    1s 3ms/step -  
 accuracy: 0.5483 - loss: 1.1836 - val\_accuracy: 0.5634 - val\_loss: 1.1478  
 Epoch 18/50  
 329/329                    1s 3ms/step -  
 accuracy: 0.5501 - loss: 1.1745 - val\_accuracy: 0.5619 - val\_loss: 1.1464  
 Epoch 19/50  
 329/329                    1s 3ms/step -  
 accuracy: 0.5529 - loss: 1.1725 - val\_accuracy: 0.5664 - val\_loss: 1.1402  
 Epoch 20/50  
 329/329                    1s 3ms/step -  
 accuracy: 0.5531 - loss: 1.1853 - val\_accuracy: 0.5668 - val\_loss: 1.1439  
 Epoch 21/50  
 329/329                    2s 4ms/step -  
 accuracy: 0.5511 - loss: 1.1853 - val\_accuracy: 0.5613 - val\_loss: 1.1364  
 Epoch 22/50  
 329/329                    1s 4ms/step -  
 accuracy: 0.5591 - loss: 1.1629 - val\_accuracy: 0.5668 - val\_loss: 1.1399  
 Epoch 23/50  
 329/329                    2s 3ms/step -  
 accuracy: 0.5608 - loss: 1.1638 - val\_accuracy: 0.5649 - val\_loss: 1.1331  
 Epoch 24/50  
 329/329                    1s 3ms/step -  
 accuracy: 0.5642 - loss: 1.1627 - val\_accuracy: 0.5649 - val\_loss: 1.1336  
 Epoch 25/50  
 329/329                    1s 3ms/step -  
 accuracy: 0.5619 - loss: 1.1544 - val\_accuracy: 0.5683 - val\_loss: 1.1359  
 Epoch 26/50  
 329/329                    1s 3ms/step -  
 accuracy: 0.5592 - loss: 1.1644 - val\_accuracy: 0.5649 - val\_loss: 1.1355  
 Epoch 27/50  
 329/329                    1s 3ms/step -  
 accuracy: 0.5655 - loss: 1.1546 - val\_accuracy: 0.5685 - val\_loss: 1.1348  
 Epoch 28/50  
 329/329                    1s 3ms/step -  
 accuracy: 0.5582 - loss: 1.1572 - val\_accuracy: 0.5653 - val\_loss: 1.1334  
 Epoch 29/50  
 329/329                    1s 3ms/step -  
 accuracy: 0.5634 - loss: 1.1465 - val\_accuracy: 0.5708 - val\_loss: 1.1299  
 Epoch 30/50

329/329                    1s 3ms/step -  
 accuracy: 0.5592 - loss: 1.1551 - val\_accuracy: 0.5697 - val\_loss: 1.1307  
 Epoch 31/50  
 329/329                    1s 4ms/step -  
 accuracy: 0.5651 - loss: 1.1452 - val\_accuracy: 0.5655 - val\_loss: 1.1298  
 Epoch 32/50  
 329/329                    1s 4ms/step -  
 accuracy: 0.5672 - loss: 1.1495 - val\_accuracy: 0.5676 - val\_loss: 1.1288  
 Epoch 33/50  
 329/329                    2s 3ms/step -  
 accuracy: 0.5653 - loss: 1.1581 - val\_accuracy: 0.5672 - val\_loss: 1.1330  
 Epoch 34/50  
 329/329                    1s 3ms/step -  
 accuracy: 0.5631 - loss: 1.1513 - val\_accuracy: 0.5666 - val\_loss: 1.1301  
 Epoch 35/50  
 329/329                    1s 3ms/step -  
 accuracy: 0.5574 - loss: 1.1558 - val\_accuracy: 0.5649 - val\_loss: 1.1287  
 Epoch 36/50  
 329/329                    1s 3ms/step -  
 accuracy: 0.5657 - loss: 1.1498 - val\_accuracy: 0.5668 - val\_loss: 1.1279  
 Epoch 37/50  
 329/329                    1s 3ms/step -  
 accuracy: 0.5659 - loss: 1.1372 - val\_accuracy: 0.5641 - val\_loss: 1.1285  
 Epoch 38/50  
 329/329                    1s 3ms/step -  
 accuracy: 0.5701 - loss: 1.1474 - val\_accuracy: 0.5689 - val\_loss: 1.1272  
 Epoch 39/50  
 329/329                    1s 3ms/step -  
 accuracy: 0.5590 - loss: 1.1509 - val\_accuracy: 0.5706 - val\_loss: 1.1248  
 Epoch 40/50  
 329/329                    1s 3ms/step -  
 accuracy: 0.5577 - loss: 1.1445 - val\_accuracy: 0.5685 - val\_loss: 1.1272  
 Epoch 41/50  
 329/329                    1s 3ms/step -  
 accuracy: 0.5662 - loss: 1.1495 - val\_accuracy: 0.5727 - val\_loss: 1.1268  
 Epoch 42/50  
 329/329                    2s 4ms/step -  
 accuracy: 0.5671 - loss: 1.1349 - val\_accuracy: 0.5693 - val\_loss: 1.1243  
 Epoch 43/50  
 329/329                    1s 4ms/step -  
 accuracy: 0.5704 - loss: 1.1349 - val\_accuracy: 0.5660 - val\_loss: 1.1264  
 Epoch 44/50  
 329/329                    2s 3ms/step -  
 accuracy: 0.5715 - loss: 1.1448 - val\_accuracy: 0.5714 - val\_loss: 1.1249  
 Epoch 45/50  
 329/329                    1s 3ms/step -  
 accuracy: 0.5670 - loss: 1.1305 - val\_accuracy: 0.5706 - val\_loss: 1.1290  
 Epoch 46/50

```

329/329          1s 3ms/step -
accuracy: 0.5779 - loss: 1.1321 - val_accuracy: 0.5721 - val_loss: 1.1230
Epoch 47/50
329/329          1s 3ms/step -
accuracy: 0.5666 - loss: 1.1420 - val_accuracy: 0.5683 - val_loss: 1.1264
Epoch 48/50
329/329          1s 3ms/step -
accuracy: 0.5656 - loss: 1.1372 - val_accuracy: 0.5664 - val_loss: 1.1226
Epoch 49/50
329/329          1s 3ms/step -
accuracy: 0.5757 - loss: 1.1244 - val_accuracy: 0.5699 - val_loss: 1.1248
Epoch 50/50
329/329          1s 3ms/step -
accuracy: 0.5687 - loss: 1.1366 - val_accuracy: 0.5666 - val_loss: 1.1280

```

## 0.6 Evaluate the deep learning model

### 0.6.1 Subtask:

Add code to evaluate the trained deep learning model on the test set and print out relevant metrics like accuracy, F1-score, and a classification report.

**Reasoning:** Evaluate the trained deep learning model on the scaled test data to assess its performance using accuracy, weighted F1-score, and a classification report.

```

[26]: from sklearn.metrics import accuracy_score, f1_score, classification_report
import numpy as np

# Evaluate the model on the test data
loss, accuracy = dl_model.evaluate(X_test_scaled, y_test, verbose=0)
print(f"Deep Learning Model Test Accuracy: {accuracy:.4f}")

# Make predictions and calculate F1-score and classification report
y_pred_dl_proba = dl_model.predict(X_test_scaled)
y_pred_dl = np.argmax(y_pred_dl_proba, axis=1)

f1_dl = f1_score(y_test, y_pred_dl, average='weighted')
class_report_dl = classification_report(y_test, y_pred_dl,
    target_names=genre_encoder.classes_)

print(f"Deep Learning Model Weighted F1 Score: {f1_dl:.4f}")
print("\nDeep Learning Model Classification Report:")
print(class_report_dl)

```

```

Deep Learning Model Test Accuracy: 0.5651
206/206          0s 1ms/step
Deep Learning Model Weighted F1 Score: 0.5585

```

Deep Learning Model Classification Report:

	precision	recall	f1-score	support
edm	0.69	0.67	0.68	1209
latin	0.54	0.40	0.46	1031
pop	0.40	0.42	0.41	1102
r&b	0.54	0.40	0.46	1086
rap	0.54	0.71	0.61	1149
rock	0.67	0.78	0.72	990
accuracy			0.57	6567
macro avg	0.56	0.56	0.56	6567
weighted avg	0.56	0.57	0.56	6567

## 0.7 Compare models

### 0.7.1 Subtask:

Add a markdown section to compare the performance of the deep learning model with the traditional machine learning models you've already evaluated.

**Reasoning:** Create a markdown section to compare the performance of the deep learning model with the traditional machine learning models using the evaluation metrics obtained.

## 0.8 Visualize deep learning results

### 0.8.1 Subtask:

Generate code to visualize the performance of the deep learning model, such as a confusion matrix or training history plots.

**Reasoning:** Generate code to visualize the deep learning model's performance using a confusion matrix and training history plots to gain insights into its predictions and training progress.

```
[51]: import os
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix
import numpy as np

# Ensure the directory exists
if not os.path.exists('images/models'):
    os.makedirs('images/models')

# Set up white background and golden color scheme
plt.style.use('default')
golden_palette = ['#8B4513', '#CD853F', '#D4AF37', '#DAA520', '#B8860B',
                  ↪ '#A0522D']

# Confusion matrix for Deep Learning Model
```

```

plt.figure(figsize=(12, 10))
cm_dl = confusion_matrix(y_test, y_pred_dl)
sns.heatmap(cm_dl, annot=True, fmt='d', cmap=golden_palette,
             xticklabels=genre_encoder.classes_,
             yticklabels=genre_encoder.classes_)
plt.title('Deep Learning Model Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.xticks(rotation=45)
plt.yticks(rotation=0)
plt.tight_layout()
plt.savefig('images/models/dl_confusion_matrix.png')
plt.show()




# Plot training history (accuracy and loss)
plt.figure(figsize=(12, 6))

# Plot accuracy
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Deep Learning Model Training and Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.grid(alpha=0.3)

# Plot loss
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Deep Learning Model Training and Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.grid(alpha=0.3)

plt.tight_layout()
plt.savefig('images/models/dl_training_history.png')
plt.show()

```



# 1 Task

Integrate PCA into the model development process by applying it before training both traditional machine learning models (like XGBoost) and a deep learning model. Evaluate the performance of all models (trained on original and PCA-transformed data), compare their results, and visualize the outcomes using the specified plotting style. Finally, summarize the findings and conclude the project.

## 1.1 Relocate and modify pca

### 1.1.1 Subtask:

Move the PCA code to a new section before model development and modify it to return the PCA-transformed training and test data.

**Reasoning:** Create a new markdown cell titled “Dimensionality Reduction (PCA)” before the “Model Development” section to introduce the PCA step.

```
[28]: # Create a new markdown cell with the specified title
```

**Reasoning:** Copy the existing PCA code, modify it to perform PCA on the scaled training and test data, and store the results in new variables. Remove the plotting code as per the instructions.

```
[29]: from sklearn.decomposition import PCA

# Apply PCA for dimensionality reduction
pca = PCA(n_components=0.95) # Retain 95% of variance
X_train_pca = pca.fit_transform(X_train_scaled)
X_test_pca = pca.transform(X_test_scaled)

print(f"Original number of features: {X_train_scaled.shape[1]}")
print(f"Reduced number of features after PCA: {X_train_pca.shape[1]}")
```

Original number of features: 14

Reduced number of features after PCA: 12

## 1.2 Train traditional ml models with pca

### 1.2.1 Subtask:

Train the selected traditional machine learning models (like XGBoost) using the PCA-transformed training data.

**Reasoning:** Train the traditional machine learning models using the PCA-transformed training data.

```
[30]: # Train traditional models on PCA-transformed data
models_pca = {}
```



```

for name, model in models.items():
    start_time = time.time()
    model.fit(X_train_pca, y_train)
    models_pca[f'{name}_PCA'] = model
    print(f"{name} (PCA) - Training Time: {time.time() - start_time:.2f}s")

```

Random Forest (PCA) - Training Time: 15.89s

Gradient Boosting (PCA) - Training Time: 114.66s

/usr/local/lib/python3.11/dist-packages/xgboost/training.py:183: UserWarning:

[02:26:21] WARNING: /workspace/src/learner.cc:738:

Parameters: { "use\_label\_encoder" } are not used.

```
bst.update(dtrain, iteration=i, fobj=obj)
```

XGBoost (PCA) - Training Time: 5.36s

SVM (PCA) - Training Time: 242.56s

**Reasoning:** Evaluate the traditional machine learning models trained on PCA-transformed data using the PCA-transformed test data.

```

[31]: from sklearn.metrics import accuracy_score, f1_score, classification_report

# Evaluate traditional models on PCA-transformed test data
results_pca = {}
for name, model in models_pca.items():
    y_pred_pca = model.predict(X_test_pca)
    accuracy_pca = accuracy_score(y_test, y_pred_pca)
    f1_pca = f1_score(y_test, y_pred_pca, average='weighted')
    results_pca[name] = {
        'test_accuracy': accuracy_pca,
        'weighted_f1': f1_pca
    }
    print(f"\n{name} - Test Accuracy: {accuracy_pca:.4f}")
    print(f"{name} - Weighted F1 Score: {f1_pca:.4f}")
    # Print classification report for detailed analysis
    # print(f"\n{name} Classification Report:\n", classification_report(y_test,
    ↪ y_pred_pca, target_names=genre_encoder.classes_))

```

Random Forest\_PCA - Test Accuracy: 0.5167

Random Forest\_PCA - Weighted F1 Score: 0.5127

Gradient Boosting\_PCA - Test Accuracy: 0.5071

Gradient Boosting\_PCA - Weighted F1 Score: 0.5055

XGBoost\_PCA - Test Accuracy: 0.5187

XGBoost\_PCA - Weighted F1 Score: 0.5159

SVM\_PCA - Test Accuracy: 0.5381

SVM\_PCA - Weighted F1 Score: 0.5350

### 1.3 Define and compile deep learning model for pca data

#### 1.3.1 Subtask:

Define and compile a deep learning model suitable for the dimensionality of the PCA-transformed data.

**Reasoning:** Define a deep learning model architecture for music genre classification using Keras Sequential model, dense layers with ReLU activation, dropout layers, and a final dense layer with softmax activation, considering the PCA-transformed input shape. Compile the model with the specified optimizer, loss function, and metrics.

```
[32]: from tensorflow.keras.models import Sequential
      from tensorflow.keras.layers import Dense, Dropout
      import numpy as np

      # Determine the number of unique genres
      num_genres = len(np.unique(y_train))

      # Define the deep learning model for PCA-transformed data
      dl_model_pca = Sequential()

      # Input layer and first hidden layer with dropout, using PCA-transformed input_
      ↪shape
      dl_model_pca.add(Dense(128, activation='relu', input_shape=(X_train_pca.
      ↪shape[1],)))
      dl_model_pca.add(Dropout(0.3)) # Added dropout

      # Second hidden layer with dropout
      dl_model_pca.add(Dense(64, activation='relu'))
      dl_model_pca.add(Dropout(0.3)) # Added dropout

      # Output layer
      dl_model_pca.add(Dense(num_genres, activation='softmax'))

      # Compile the model
      dl_model_pca.compile(optimizer='adam', loss='sparse_categorical_crossentropy', ↪
      ↪metrics=['accuracy']) # Added compile

      dl_model_pca.summary()
```

```
/usr/local/lib/python3.11/dist-packages/keras/src/layers/core/dense.py:93:
UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When
using Sequential models, prefer using an `Input(shape)` object as the first
layer in the model instead.
```

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
dense_3 (Dense)	(None, 128)	1,664
dropout_2 (Dropout)	(None, 128)	0
dense_4 (Dense)	(None, 64)	8,256
dropout_3 (Dropout)	(None, 64)	0
dense_5 (Dense)	(None, 6)	390

Total params: 10,310 (40.27 KB)

Trainable params: 10,310 (40.27 KB)

Non-trainable params: 0 (0.00 B)

## 1.4 Train deep learning model with pca

### 1.4.1 Subtask:

Train the deep learning model using the PCA-transformed training data.

**Reasoning:** Train the deep learning model on the PCA-transformed training data with specified epochs, batch size, and validation split, including the EarlyStopping callback.

```
[33]: from tensorflow.keras.callbacks import EarlyStopping

# Instantiate EarlyStopping callback
early_stopping = EarlyStopping(monitor='val_loss', patience=15,
                                restore_best_weights=True)

# Train the deep learning model for PCA-transformed data
history_pca = dl_model_pca.fit(X_train_pca, y_train,
                                epochs=50,
                                batch_size=64,
                                validation_split=0.2,
                                callbacks=[early_stopping])
```

Epoch 1/50

329/329

2s 3ms/step -

accuracy: 0.3734 - loss: 1.5598 - val\_accuracy: 0.4931 - val\_loss: 1.3119  
 Epoch 2/50  
 329/329 2s 4ms/step -  
 accuracy: 0.4690 - loss: 1.3641 - val\_accuracy: 0.5069 - val\_loss: 1.2770  
 Epoch 3/50  
 329/329 2s 3ms/step -  
 accuracy: 0.4896 - loss: 1.3332 - val\_accuracy: 0.5147 - val\_loss: 1.2530  
 Epoch 4/50  
 329/329 1s 3ms/step -  
 accuracy: 0.4939 - loss: 1.2996 - val\_accuracy: 0.5179 - val\_loss: 1.2427  
 Epoch 5/50  
 329/329 1s 3ms/step -  
 accuracy: 0.5125 - loss: 1.2768 - val\_accuracy: 0.5232 - val\_loss: 1.2325  
 Epoch 6/50  
 329/329 1s 3ms/step -  
 accuracy: 0.5048 - loss: 1.2784 - val\_accuracy: 0.5282 - val\_loss: 1.2199  
 Epoch 7/50  
 329/329 1s 3ms/step -  
 accuracy: 0.5146 - loss: 1.2562 - val\_accuracy: 0.5344 - val\_loss: 1.2181  
 Epoch 8/50  
 329/329 1s 3ms/step -  
 accuracy: 0.5151 - loss: 1.2623 - val\_accuracy: 0.5339 - val\_loss: 1.2101  
 Epoch 9/50  
 329/329 1s 3ms/step -  
 accuracy: 0.5162 - loss: 1.2502 - val\_accuracy: 0.5398 - val\_loss: 1.2044  
 Epoch 10/50  
 329/329 1s 3ms/step -  
 accuracy: 0.5267 - loss: 1.2252 - val\_accuracy: 0.5375 - val\_loss: 1.1959  
 Epoch 11/50  
 329/329 1s 3ms/step -  
 accuracy: 0.5310 - loss: 1.2379 - val\_accuracy: 0.5411 - val\_loss: 1.1916  
 Epoch 12/50  
 329/329 1s 3ms/step -  
 accuracy: 0.5333 - loss: 1.2231 - val\_accuracy: 0.5394 - val\_loss: 1.1930  
 Epoch 13/50  
 329/329 1s 4ms/step -  
 accuracy: 0.5334 - loss: 1.2193 - val\_accuracy: 0.5426 - val\_loss: 1.1885  
 Epoch 14/50  
 329/329 1s 4ms/step -  
 accuracy: 0.5291 - loss: 1.2257 - val\_accuracy: 0.5434 - val\_loss: 1.1827  
 Epoch 15/50  
 329/329 2s 3ms/step -  
 accuracy: 0.5409 - loss: 1.2067 - val\_accuracy: 0.5411 - val\_loss: 1.1828  
 Epoch 16/50  
 329/329 1s 3ms/step -  
 accuracy: 0.5452 - loss: 1.1990 - val\_accuracy: 0.5436 - val\_loss: 1.1846  
 Epoch 17/50  
 329/329 1s 3ms/step -

accuracy: 0.5365 - loss: 1.2168 - val\_accuracy: 0.5463 - val\_loss: 1.1785  
 Epoch 18/50  
 329/329 1s 3ms/step -  
 accuracy: 0.5437 - loss: 1.2000 - val\_accuracy: 0.5489 - val\_loss: 1.1716  
 Epoch 19/50  
 329/329 1s 3ms/step -  
 accuracy: 0.5423 - loss: 1.1983 - val\_accuracy: 0.5493 - val\_loss: 1.1729  
 Epoch 20/50  
 329/329 1s 3ms/step -  
 accuracy: 0.5404 - loss: 1.1911 - val\_accuracy: 0.5451 - val\_loss: 1.1711  
 Epoch 21/50  
 329/329 1s 3ms/step -  
 accuracy: 0.5406 - loss: 1.1913 - val\_accuracy: 0.5472 - val\_loss: 1.1715  
 Epoch 22/50  
 329/329 1s 3ms/step -  
 accuracy: 0.5419 - loss: 1.2085 - val\_accuracy: 0.5482 - val\_loss: 1.1702  
 Epoch 23/50  
 329/329 2s 4ms/step -  
 accuracy: 0.5394 - loss: 1.1978 - val\_accuracy: 0.5522 - val\_loss: 1.1681  
 Epoch 24/50  
 329/329 1s 4ms/step -  
 accuracy: 0.5442 - loss: 1.1926 - val\_accuracy: 0.5501 - val\_loss: 1.1676  
 Epoch 25/50  
 329/329 2s 3ms/step -  
 accuracy: 0.5472 - loss: 1.1845 - val\_accuracy: 0.5489 - val\_loss: 1.1654  
 Epoch 26/50  
 329/329 1s 3ms/step -  
 accuracy: 0.5376 - loss: 1.1932 - val\_accuracy: 0.5508 - val\_loss: 1.1612  
 Epoch 27/50  
 329/329 1s 3ms/step -  
 accuracy: 0.5429 - loss: 1.1932 - val\_accuracy: 0.5464 - val\_loss: 1.1680  
 Epoch 28/50  
 329/329 1s 3ms/step -  
 accuracy: 0.5458 - loss: 1.1908 - val\_accuracy: 0.5508 - val\_loss: 1.1611  
 Epoch 29/50  
 329/329 1s 3ms/step -  
 accuracy: 0.5479 - loss: 1.1852 - val\_accuracy: 0.5546 - val\_loss: 1.1609  
 Epoch 30/50  
 329/329 1s 3ms/step -  
 accuracy: 0.5480 - loss: 1.1800 - val\_accuracy: 0.5546 - val\_loss: 1.1594  
 Epoch 31/50  
 329/329 1s 3ms/step -  
 accuracy: 0.5445 - loss: 1.1845 - val\_accuracy: 0.5537 - val\_loss: 1.1603  
 Epoch 32/50  
 329/329 1s 3ms/step -  
 accuracy: 0.5528 - loss: 1.1721 - val\_accuracy: 0.5537 - val\_loss: 1.1583  
 Epoch 33/50  
 329/329 1s 3ms/step -

accuracy: 0.5500 - loss: 1.1715 - val\_accuracy: 0.5584 - val\_loss: 1.1563  
 Epoch 34/50  
 329/329 2s 4ms/step -  
 accuracy: 0.5464 - loss: 1.1859 - val\_accuracy: 0.5522 - val\_loss: 1.1598  
 Epoch 35/50  
 329/329 2s 5ms/step -  
 accuracy: 0.5426 - loss: 1.1893 - val\_accuracy: 0.5569 - val\_loss: 1.1568  
 Epoch 36/50  
 329/329 2s 5ms/step -  
 accuracy: 0.5470 - loss: 1.1834 - val\_accuracy: 0.5546 - val\_loss: 1.1549  
 Epoch 37/50  
 329/329 2s 5ms/step -  
 accuracy: 0.5606 - loss: 1.1667 - val\_accuracy: 0.5518 - val\_loss: 1.1548  
 Epoch 38/50  
 329/329 1s 4ms/step -  
 accuracy: 0.5538 - loss: 1.1724 - val\_accuracy: 0.5571 - val\_loss: 1.1562  
 Epoch 39/50  
 329/329 1s 3ms/step -  
 accuracy: 0.5517 - loss: 1.1735 - val\_accuracy: 0.5569 - val\_loss: 1.1582  
 Epoch 40/50  
 329/329 1s 3ms/step -  
 accuracy: 0.5518 - loss: 1.1690 - val\_accuracy: 0.5542 - val\_loss: 1.1559  
 Epoch 41/50  
 329/329 1s 3ms/step -  
 accuracy: 0.5610 - loss: 1.1576 - val\_accuracy: 0.5569 - val\_loss: 1.1546  
 Epoch 42/50  
 329/329 1s 3ms/step -  
 accuracy: 0.5490 - loss: 1.1767 - val\_accuracy: 0.5531 - val\_loss: 1.1571  
 Epoch 43/50  
 329/329 1s 3ms/step -  
 accuracy: 0.5493 - loss: 1.1776 - val\_accuracy: 0.5535 - val\_loss: 1.1541  
 Epoch 44/50  
 329/329 1s 3ms/step -  
 accuracy: 0.5474 - loss: 1.1762 - val\_accuracy: 0.5531 - val\_loss: 1.1530  
 Epoch 45/50  
 329/329 1s 3ms/step -  
 accuracy: 0.5531 - loss: 1.1780 - val\_accuracy: 0.5533 - val\_loss: 1.1513  
 Epoch 46/50  
 329/329 1s 3ms/step -  
 accuracy: 0.5550 - loss: 1.1620 - val\_accuracy: 0.5537 - val\_loss: 1.1512  
 Epoch 47/50  
 329/329 1s 3ms/step -  
 accuracy: 0.5472 - loss: 1.1768 - val\_accuracy: 0.5552 - val\_loss: 1.1532  
 Epoch 48/50  
 329/329 2s 4ms/step -  
 accuracy: 0.5593 - loss: 1.1589 - val\_accuracy: 0.5571 - val\_loss: 1.1518  
 Epoch 49/50  
 329/329 1s 4ms/step -

```
accuracy: 0.5519 - loss: 1.1731 - val_accuracy: 0.5552 - val_loss: 1.1516
Epoch 50/50
329/329          2s 3ms/step -
accuracy: 0.5562 - loss: 1.1733 - val_accuracy: 0.5565 - val_loss: 1.1528
```

## 1.5 Evaluate deep learning model with PCA

### 1.5.1 Subtask:

Evaluate the performance of the deep learning model trained on PCA data and compare the results with previous deep learning evaluations.

**Reasoning:** Evaluate the deep learning model trained on PCA-transformed data using the PCA-transformed test data to assess its performance with accuracy, weighted F1-score, and a classification report, and compare with previous deep learning evaluations.

```
[34]: from sklearn.metrics import accuracy_score, f1_score, classification_report
import numpy as np

# Evaluate the deep learning model trained on PCA data
loss_dl_pca, accuracy_dl_pca = dl_model_pca.evaluate(X_test_pca, y_test,
    verbose=0)
print(f"Deep Learning Model (PCA) Test Accuracy: {accuracy_dl_pca:.4f}")

# Make predictions and calculate F1-score and classification report for the DL
    model with PCA
y_pred_dl_pca_proba = dl_model_pca.predict(X_test_pca)
y_pred_dl_pca = np.argmax(y_pred_dl_pca_proba, axis=1)

f1_dl_pca = f1_score(y_test, y_pred_dl_pca, average='weighted')
class_report_dl_pca = classification_report(y_test, y_pred_dl_pca,
    target_names=genre_encoder.classes_)

print(f"Deep Learning Model (PCA) Weighted F1 Score: {f1_dl_pca:.4f}")
print("\nDeep Learning Model (PCA) Classification Report:")
print(class_report_dl_pca)
```

```
Deep Learning Model (PCA) Test Accuracy: 0.5555
206/206          0s 1ms/step
Deep Learning Model (PCA) Weighted F1 Score: 0.5509
```

```
Deep Learning Model (PCA) Classification Report:
      precision    recall  f1-score   support

edm          0.67       0.66       0.67       1209
latin         0.50       0.42       0.46       1031
pop           0.39       0.39       0.39       1102
r&b           0.52       0.43       0.47       1086
rap           0.56       0.68       0.62       1149
```

rock	0.66	0.73	0.70	990
accuracy			0.56	6567
macro avg	0.55	0.55	0.55	6567
weighted avg	0.55	0.56	0.55	6567

## 1.6 Compare all model performances

### 1.6.1 Subtask:

Create a comprehensive comparison of all model performances (traditional ML on original data, traditional ML on PCA data, deep learning on original data, deep learning on PCA data).

**Reasoning:** Create a markdown section to compare the performance of all trained models, including traditional ML models on original and PCA data, and deep learning models on original and PCA data, using their respective evaluation metrics.

```
[35]: import pandas as pd

# Create a dictionary to store the results
all_results = {
    'Traditional ML (Original)': results,
    'Traditional ML (PCA)': results_pca,
    'Deep Learning (Original)': {'Test Accuracy': accuracy, 'Weighted F1 Score':
    ↪ f1},
    'Deep Learning (PCA)': {'Test Accuracy': accuracy_dl_pca, 'Weighted F1_
    ↪ Score': f1_dl_pca}
}

# Convert results to a pandas DataFrame for easy comparison
comparison_data = []

for model_type, model_results in all_results.items():
    if 'Traditional ML' in model_type:
        for model_name, metrics in model_results.items():
            comparison_data.append({
                'Model Type': model_type,
                'Model Name': model_name,
                'Test Accuracy': metrics.get('test_accuracy', metrics.
    ↪ get('cv_accuracy')), # Handle both keys
                'Weighted F1 Score': metrics.get('weighted_f1')
            })
    else:
        comparison_data.append({
            'Model Type': model_type,
            'Model Name': model_type, # Use model type as name for DL
            'Test Accuracy': model_results['Test Accuracy'],
            'Weighted F1 Score': model_results['Weighted F1 Score']
        })
```



```

    })

comparison_df = pd.DataFrame(comparison_data)

# Display the comparison table
print("Model Performance Comparison:")
display(comparison_df.round(4))

```

Model Performance Comparison:

	Model Type	Model Name	Test Accuracy \
0	Traditional ML (Original)	Random Forest	0.5747
1	Traditional ML (Original)	Gradient Boosting	0.5715
2	Traditional ML (Original)	XGBoost	0.5876
3	Traditional ML (Original)	SVM	0.5548
4	Traditional ML (PCA)	Random Forest_PCA	0.5167
5	Traditional ML (PCA)	Gradient Boosting_PCA	0.5071
6	Traditional ML (PCA)	XGBoost_PCA	0.5187
7	Traditional ML (PCA)	SVM_PCA	0.5381
8	Deep Learning (Original)	Deep Learning (Original)	0.5651
9	Deep Learning (PCA)	Deep Learning (PCA)	0.5555

	Weighted F1 Score
0	NaN
1	NaN
2	NaN
3	NaN
4	0.5127
5	0.5055
6	0.5159
7	0.5350
8	0.5940
9	0.5509

## 2 Task

Use AutoML (AutoGluon) to improve the model accuracy and compare its performance with previously trained models.

### 2.1 Install autogluon

#### 2.1.1 Subtask:

Add code to install the AutoGluon library.

**Reasoning:** Install the AutoGluon library using pip.

```
[36]: # %pip install --quiet autogluon
```

## 2.2 Train AutoGluon Model

### 2.2.1 Subtask:

Add code to train an AutoGluon TabularPredictor on your training data for music genre classification.

**Reasoning:** Train an AutoGluon TabularPredictor on the training data for music genre classification, specifying the label column and evaluation metric.

```
[37]: from autogluon.tabular import TabularPredictor
import time

# Prepare data for AutoGluon
# AutoGluon works directly with pandas DataFrames, so we can use the original
↳ X_train and y_train
# We need to combine X_train and y_train into a single DataFrame for AutoGluon
train_data = X_train.copy()
train_data['genre_encoded'] = y_train

# Specify the label column
label = 'genre_encoded'

# Initialize and train the AutoGluon predictor
start_time = time.time()
# Exclude NeuralNetTorch model
predictor = TabularPredictor(label=label, eval_metric='accuracy').
↳ fit(train_data, excluded_model_types=['NN_TORCH'])
autogluon_train_time = time.time() - start_time

print(f"AutoGluon Model Training Time: {autogluon_train_time:.2f}s")
```

No path specified. Models will be saved in: "AutogluonModels/ag-20250807\_023147"

Verbosity: 2 (Standard Logging)

===== System Info =====

AutoGluon Version: 1.4.0

Python Version: 3.11.13

Operating System: Linux

Platform Machine: x86\_64

Platform Version: #1 SMP PREEMPT\_DYNAMIC Sun Mar 30 16:01:29 UTC 2025

CPU Count: 2

Memory Avail: 10.17 GB / 12.67 GB (80.3%)

Disk Space Avail: 56.75 GB / 107.72 GB (52.7%)

=====

No presets specified! To achieve strong results with AutoGluon, it is recommended to use the available presets. Defaulting to `medium`...

Recommended Presets (For more details refer to

<https://auto.gluon.ai/stable/tutorials/tabular/tabular-essentials.html#presets>):

presets='extreme' : New in v1.4: Massively better than 'best' on

datasets <30000 samples by using new models meta-learned on <https://tabarena.ai>: TabPFNv2, TabICL, Mitra, and TabM. Absolute best accuracy. Requires a GPU. Recommended 64 GB CPU memory and 32+ GB GPU memory.

presets='best' : Maximize accuracy. Recommended for most users. Use in competitions and benchmarks.

presets='high' : Strong accuracy with fast inference speed.

presets='good' : Good accuracy with very fast inference speed.

presets='medium' : Fast training time, ideal for initial prototyping.

Using hyperparameters preset: hyperparameters='default'

Beginning AutoGluon training ...

AutoGluon will save models to "/content/AutogluonModels/ag-20250807\_023147"

Train Data Rows: 26266

Train Data Columns: 14

Label Column: genre\_encoded

AutoGluon infers your prediction problem is: 'multiclass' (because dtype of label-column == int, but few unique label-values observed).

6 unique label values: [1, 2, 3, 5, 0, 4]

If 'multiclass' is not the correct problem\_type, please manually specify the problem\_type parameter during Predictor init (You may specify problem\_type as one of: ['binary', 'multiclass', 'regression', 'quantile'])

Problem Type: multiclass

Preprocessing data ...

Train Data Class Count: 6

Using Feature Generators to preprocess the data ...

Fitting AutoMLPipelineFeatureGenerator...

Available Memory: 10414.84 MB

Train Data (Original) Memory Usage: 2.81 MB (0.0% of available memory)

Inferring data type of each feature based on column values. Set

feature\_metadata\_in to manually specify special dtypes of the features.

Stage 1 Generators:

Fitting AsTypeFeatureGenerator...

Note: Converting 1 features to boolean dtype as they only contain 2 unique values.

Stage 2 Generators:

Fitting FillNaFeatureGenerator...

Stage 3 Generators:

Fitting IdentityFeatureGenerator...

Stage 4 Generators:

Fitting DropUniqueFeatureGenerator...

Stage 5 Generators:

Fitting DropDuplicatesFeatureGenerator...

Types of features in original data (raw dtype, special dtypes):

('float', []) : 11 | ['danceability', 'energy', 'loudness', 'speechiness', 'acousticness', ...]

('int', []) : 3 | ['key', 'mode', 'track\_popularity']

Types of features in processed data (raw dtype, special dtypes):

('float', []) : 11 | ['danceability', 'energy', 'loudness', 'speechiness', 'acousticness', ...]

```

        ('int', [])          : 2 | ['key', 'track_popularity']
        ('int', ['bool'])    : 1 | ['mode']
0.1s = Fit runtime
14 features in original data used to generate 14 features in processed
data.
Train Data (Processed) Memory Usage: 2.63 MB (0.0% of available memory)
Data preprocessing and feature engineering runtime = 0.15s ...
AutoGluon will gauge predictive performance using evaluation metric: 'accuracy'
To change this, specify the eval_metric parameter of Predictor()
Automatically generating train/validation split with
holdout_frac=0.09518008071270845, Train Rows: 23766, Val Rows: 2500
User-specified model hyperparameters to be fit:
{
    'NN_TORCH': [{}],
    'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}, {}],
{'learning_rate': 0.03, 'num_leaves': 128, 'feature_fraction': 0.9,
'min_data_in_leaf': 3, 'ag_args': {'name_suffix': 'Large', 'priority': 0,
'hyperparameter_tune_kwargs': None}}],
    'CAT': [{}],
    'XGB': [{}],
    'FASTAI': [{}],
    'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
    'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
}
Excluded models: ['NN_TORCH'] (Specified by `excluded_model_types`)
Fitting 10 L1 models, fit_strategy="sequential" ...
Fitting model: NeuralNetFastAI ...
Fitting with cpus=1, gpus=0, mem=0.0/10.0 GB
0.5836 = Validation score (accuracy)
29.1s = Training runtime
0.03s = Validation runtime
Fitting model: LightGBMXT ...
Fitting with cpus=1, gpus=0, mem=0.0/9.9 GB
0.588 = Validation score (accuracy)
12.41s = Training runtime
0.67s = Validation runtime
Fitting model: LightGBM ...
Fitting with cpus=1, gpus=0, mem=0.0/9.9 GB
0.6024 = Validation score (accuracy)
14.19s = Training runtime

```

```

    0.84s      = Validation runtime
Fitting model: RandomForestGini ...
    Fitting with cpus=2, gpus=0, mem=0.1/9.9 GB
    0.5772    = Validation score    (accuracy)
    26.68s    = Training    runtime
    0.55s     = Validation runtime
Fitting model: RandomForestEntr ...
    Fitting with cpus=2, gpus=0, mem=0.1/9.7 GB
    0.5808    = Validation score    (accuracy)
    42.37s    = Training    runtime
    0.27s     = Validation runtime
Fitting model: CatBoost ...
    Fitting with cpus=1, gpus=0
    0.5712    = Validation score    (accuracy)
    10.8s     = Training    runtime
    0.01s     = Validation runtime
Fitting model: ExtraTreesGini ...
    Fitting with cpus=2, gpus=0, mem=0.1/9.6 GB
    0.5632    = Validation score    (accuracy)
    10.46s    = Training    runtime
    0.29s     = Validation runtime
Fitting model: ExtraTreesEntr ...
    Fitting with cpus=2, gpus=0, mem=0.1/9.6 GB
    0.5608    = Validation score    (accuracy)
    9.54s     = Training    runtime
    0.3s      = Validation runtime
Fitting model: XGBoost ...
    Fitting with cpus=1, gpus=0
    0.5904    = Validation score    (accuracy)
    16.21s    = Training    runtime
    0.34s     = Validation runtime
Fitting model: LightGBMLarge ...
    Fitting with cpus=1, gpus=0, mem=0.1/9.6 GB
    0.5936    = Validation score    (accuracy)
    16.83s    = Training    runtime
    0.6s      = Validation runtime
Fitting model: WeightedEnsemble_L2 ...
    Ensemble Weights: {'LightGBM': 1.0}
    0.6024    = Validation score    (accuracy)
    0.18s     = Training    runtime
    0.0s      = Validation runtime
AutoGluon training complete, total runtime = 220.37s ... Best model:
WeightedEnsemble_L2 | Estimated inference throughput: 2979.7 rows/s (2500 batch
size)
TabularPredictor saved. To load, use: predictor =
TabularPredictor.load("/content/AutogluonModels/ag-20250807_023147")

AutoGluon Model Training Time: 221.53s

```

## 2.3 Evaluate AutoGluon Model

### 2.3.1 Subtask:

Add code to evaluate the trained AutoGluon model on the test data and print out relevant metrics like accuracy and F1-score.

**Reasoning:** Evaluate the trained AutoGluon model on the test data to assess its performance using accuracy and weighted F1-score.

```
[38]: # Evaluate the AutoGluon model on the test data
# AutoGluon's predictor object can directly predict on a DataFrame
test_data = X_test.copy()
y_pred_autogluon = predictor.predict(test_data)
y_proba_autogluon = predictor.predict_proba(test_data)

# Calculate evaluation metrics
accuracy_autogluon = accuracy_score(y_test, y_pred_autogluon)
f1_autogluon = f1_score(y_test, y_pred_autogluon, average='weighted')

print("\n" + "="*50)
print("AutoGluon Model Performance:")
print(f"Test Accuracy: {accuracy_autogluon:.4f}")
print(f"Weighted F1 Score: {f1_autogluon:.4f}")
print("\nClassification Report:")
print(classification_report(y_test, y_pred_autogluon,
    ↪target_names=genre_encoder.classes_))
```

=====

AutoGluon Model Performance:

Test Accuracy: 0.5838

Weighted F1 Score: 0.5825

Classification Report:

	precision	recall	f1-score	support
edm	0.69	0.69	0.69	1209
latin	0.53	0.47	0.50	1031
pop	0.39	0.42	0.41	1102
r&b	0.52	0.47	0.49	1086
rap	0.61	0.67	0.64	1149
rock	0.75	0.77	0.76	990
accuracy			0.58	6567
macro avg	0.58	0.58	0.58	6567
weighted avg	0.58	0.58	0.58	6567

## 2.4 Compare all model performances

### 2.4.1 Subtask:

Create a comprehensive comparison of all model performances (traditional ML on original data, traditional ML on PCA data, deep learning on original data, deep learning on PCA data, and AutoGluon).

**Reasoning:** Create a comprehensive comparison of all model performances, including traditional ML models on original and PCA data, deep learning models on original and PCA data, and the AutoGluon model, using their respective evaluation metrics.

```
[39]: import pandas as pd

# Initialize results and results_pca as empty dictionaries if they are not
# defined
if 'results' not in locals():
    results = {}
if 'results_pca' not in locals():
    results_pca = {}

# Create a dictionary to store the results
all_results = {
    'Traditional ML (Original)': results,
    'Traditional ML (PCA)': results_pca,
}

# Add Deep Learning (Original) results if available
if 'accuracy' in locals() and 'f1' in locals():
    all_results['Deep Learning (Original)'] = {'Test Accuracy': accuracy,
    ↪ 'Weighted F1 Score': f1}

# Add Deep Learning (PCA) results if available
if 'accuracy_dl_pca' in locals() and 'f1_dl_pca' in locals():
    all_results['Deep Learning (PCA)'] = {'Test Accuracy': accuracy_dl_pca,
    ↪ 'Weighted F1 Score': f1_dl_pca}

# Add AutoGluon results if available
if 'accuracy_autogluon' in locals() and 'f1_autogluon' in locals():
    all_results['AutoGluon'] = {'Test Accuracy': accuracy_autogluon, 'Weighted
    ↪ F1 Score': f1_autogluon}

# Convert results to a pandas DataFrame for easy comparison
comparison_data = []

for model_type, model_results in all_results.items():
    if 'Traditional ML' in model_type:
        for model_name, metrics in model_results.items():
```

```

        comparison_data.append({
            'Model Type': model_type,
            'Model Name': model_name,
            'Test Accuracy': metrics.get('test_accuracy', metrics.
↪get('cv_accuracy')), # Handle both keys
            'Weighted F1 Score': metrics.get('weighted_f1')
        })
    else:
        comparison_data.append({
            'Model Type': model_type,
            'Model Name': model_type, # Use model type as name for DL and
↪AutoGluon
            'Test Accuracy': model_results['Test Accuracy'],
            'Weighted F1 Score': model_results['Weighted F1 Score']
        })

comparison_df = pd.DataFrame(comparison_data)

# Display the comparison table
print("Model Performance Comparison:")
display(comparison_df.round(4))

```

Model Performance Comparison:

	Model Type	Model Name	Test Accuracy \
0	Traditional ML (Original)	Random Forest	0.5747
1	Traditional ML (Original)	Gradient Boosting	0.5715
2	Traditional ML (Original)	XGBoost	0.5876
3	Traditional ML (Original)	SVM	0.5548
4	Traditional ML (PCA)	Random Forest_PCA	0.5167
5	Traditional ML (PCA)	Gradient Boosting_PCA	0.5071
6	Traditional ML (PCA)	XGBoost_PCA	0.5187
7	Traditional ML (PCA)	SVM_PCA	0.5381
8	Deep Learning (Original)	Deep Learning (Original)	0.5651
9	Deep Learning (PCA)	Deep Learning (PCA)	0.5555
10	AutoGluon	AutoGluon	0.5838

	Weighted F1 Score
0	NaN
1	NaN
2	NaN
3	NaN
4	0.5127
5	0.5055
6	0.5159
7	0.5350
8	0.5940



9	0.5509
10	0.5825

### 3 Task

Train an AutoGluon model on PCA-transformed data, evaluate its performance, update the comparison table and visualization to include these results, and finally summarize the findings from all models.

#### 3.1 Train autogluon with pca

Add a new section to train the AutoGluon TabularPredictor on the PCA-transformed training data (X\_train\_pca).

**Reasoning:** Prepare the PCA-transformed training data for AutoGluon and train the AutoGluon TabularPredictor on this data, measuring the training time.

```
[48]: # Apply PCA for dimensionality reduction
pca = PCA(n_components=0.95) # Retain 95% of variance
X_train_pca = pca.fit_transform(X_train_scaled)
X_test_pca = pca.transform(X_test_scaled)

print(f"Original number of features: {X_train_scaled.shape[1]}")
print(f"Reduced number of features after PCA: {X_train_pca.shape[1]}")

# Convert PCA-transformed data and target variable to DataFrames with original
# indices
X_train_pca_df = pd.DataFrame(X_train_pca, index=X_train.index)
y_train_df = pd.DataFrame(y_train, index=X_train.index,
                           columns=['genre_encoded'])

# Combine features and target
train_data_pca_combined = pd.concat([X_train_pca_df, y_train_df], axis=1)

# Drop rows with any non-finite values in the combined DataFrame
train_data_pca_cleaned = train_data_pca_combined.dropna().copy()

# Specify the label column
label = 'genre_encoded'

# Initialize and train the AutoGluon predictor on cleaned PCA data
start_time = time.time()
predictor_pca = TabularPredictor(label=label, eval_metric='accuracy').
    fit(train_data_pca_cleaned, excluded_model_types=['NN_TORCH'])
autogluon_pca_train_time = time.time() - start_time

print(f"AutoGluon Model (PCA) Training Time: {autogluon_pca_train_time:.2f}s")
```

No path specified. Models will be saved in: "AutogluonModels/ag-20250807\_025729"  
Verbosity: 2 (Standard Logging)

===== System Info =====

AutoGluon Version: 1.4.0  
Python Version: 3.11.13  
Operating System: Linux  
Platform Machine: x86\_64  
Platform Version: #1 SMP PREEMPT\_DYNAMIC Sun Mar 30 16:01:29 UTC 2025  
CPU Count: 2  
Memory Avail: 9.35 GB / 12.67 GB (73.8%)  
Disk Space Avail: 47.02 GB / 107.72 GB (43.7%)

=====

No presets specified! To achieve strong results with AutoGluon, it is recommended to use the available presets. Defaulting to `medium`...

Recommended Presets (For more details refer to <https://auto.gluon.ai/stable/tutorials/tabular/tabular-essentials.html#presets>):  
presets='extreme' : New in v1.4: Massively better than 'best' on datasets <30000 samples by using new models meta-learned on <https://tabarena.ai>: TabPFNv2, TabICL, Mitra, and TabM. Absolute best accuracy. Requires a GPU. Recommended 64 GB CPU memory and 32+ GB GPU memory.

presets='best' : Maximize accuracy. Recommended for most users. Use in competitions and benchmarks.

presets='high' : Strong accuracy with fast inference speed.

presets='good' : Good accuracy with very fast inference speed.

presets='medium' : Fast training time, ideal for initial prototyping.

Using hyperparameters preset: hyperparameters='default'

Original number of features: 14

Reduced number of features after PCA: 12

Beginning AutoGluon training ...

AutoGluon will save models to "/content/AutogluonModels/ag-20250807\_025729"

Train Data Rows: 26266

Train Data Columns: 12

Label Column: genre\_encoded

AutoGluon infers your prediction problem is: 'multiclass' (because dtype of label-column == int, but few unique label-values observed).

6 unique label values: [1, 2, 3, 5, 0, 4]

If 'multiclass' is not the correct problem\_type, please manually specify the problem\_type parameter during Predictor init (You may specify problem\_type as one of: ['binary', 'multiclass', 'regression', 'quantile'])

Problem Type: multiclass

Preprocessing data ...

Train Data Class Count: 6

Using Feature Generators to preprocess the data ...

Fitting AutoMLPipelineFeatureGenerator...

Available Memory: 9575.30 MB

Train Data (Original) Memory Usage: 2.40 MB (0.0% of available memory)

Inferring data type of each feature based on column values. Set

```

feature_metadata_in to manually specify special dtypes of the features.
Stage 1 Generators:
    Fitting AsTypeFeatureGenerator...
Stage 2 Generators:
    Fitting FillNaFeatureGenerator...
Stage 3 Generators:
    Fitting IdentityFeatureGenerator...
Stage 4 Generators:
    Fitting DropUniqueFeatureGenerator...
Stage 5 Generators:
    Fitting DropDuplicatesFeatureGenerator...
Types of features in original data (raw dtype, special dtypes):
    ('float', []) : 12 | ['0', '1', '2', '3', '4', ...]
Types of features in processed data (raw dtype, special dtypes):
    ('float', []) : 12 | ['0', '1', '2', '3', '4', ...]
0.1s = Fit runtime
12 features in original data used to generate 12 features in processed
data.
Train Data (Processed) Memory Usage: 2.40 MB (0.0% of available memory)
Data preprocessing and feature engineering runtime = 0.18s ...
AutoGluon will gauge predictive performance using evaluation metric: 'accuracy'
To change this, specify the eval_metric parameter of Predictor()
Automatically generating train/validation split with
holdout_frac=0.09518008071270845, Train Rows: 23766, Val Rows: 2500
User-specified model hyperparameters to be fit:
{
    'NN_TORCH': [{}],
    'GBM': [{'extra_trees': True, 'ag_args': {'name_suffix': 'XT'}}, {}],
{'learning_rate': 0.03, 'num_leaves': 128, 'feature_fraction': 0.9,
'min_data_in_leaf': 3, 'ag_args': {'name_suffix': 'Large', 'priority': 0,
'hyperparameter_tune_kwargs': None}}],
    'CAT': [{}],
    'XGB': [{}],
    'FASTAI': [{}],
    'RF': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
    'XT': [{'criterion': 'gini', 'ag_args': {'name_suffix': 'Gini',
'problem_types': ['binary', 'multiclass']}}, {'criterion': 'entropy', 'ag_args':
{'name_suffix': 'Entr', 'problem_types': ['binary', 'multiclass']}},
{'criterion': 'squared_error', 'ag_args': {'name_suffix': 'MSE',
'problem_types': ['regression', 'quantile']}}],
}
Excluded models: ['NN_TORCH'] (Specified by `excluded_model_types`)
Fitting 10 L1 models, fit_strategy="sequential" ...
Fitting model: NeuralNetFastAI ...

```

```

Fitting with cpus=1, gpus=0, mem=0.0/9.4 GB
0.56      = Validation score    (accuracy)
27.06s    = Training    runtime
0.03s     = Validation runtime

Fitting model: LightGBMXT ...
Fitting with cpus=1, gpus=0, mem=0.0/9.3 GB
0.5388    = Validation score    (accuracy)
8.02s     = Training    runtime
0.39s     = Validation runtime

Fitting model: LightGBM ...
Fitting with cpus=1, gpus=0, mem=0.0/9.3 GB
0.5416    = Validation score    (accuracy)
7.38s     = Training    runtime
0.38s     = Validation runtime

Fitting model: RandomForestGini ...
Fitting with cpus=2, gpus=0, mem=0.1/9.3 GB
0.5268    = Validation score    (accuracy)
39.88s    = Training    runtime
0.54s     = Validation runtime

Fitting model: RandomForestEntr ...
Fitting with cpus=2, gpus=0, mem=0.1/8.9 GB
0.522     = Validation score    (accuracy)
56.79s    = Training    runtime
0.26s     = Validation runtime

Fitting model: CatBoost ...
Fitting with cpus=1, gpus=0
0.5412    = Validation score    (accuracy)
68.9s     = Training    runtime
0.04s     = Validation runtime

Fitting model: ExtraTreesGini ...
Fitting with cpus=2, gpus=0, mem=0.1/9.1 GB
0.532     = Validation score    (accuracy)
9.07s     = Training    runtime
0.3s      = Validation runtime

Fitting model: ExtraTreesEntr ...
Fitting with cpus=2, gpus=0, mem=0.1/9.2 GB
0.5316    = Validation score    (accuracy)
10.32s    = Training    runtime
0.32s     = Validation runtime

Fitting model: XGBoost ...
Fitting with cpus=1, gpus=0
0.5296    = Validation score    (accuracy)
12.54s    = Training    runtime
0.3s      = Validation runtime

Fitting model: LightGBMLarge ...
Fitting with cpus=1, gpus=0, mem=0.1/9.2 GB
0.5368    = Validation score    (accuracy)
22.09s    = Training    runtime

```

```

0.67s      = Validation runtime
Fitting model: WeightedEnsemble_L2 ...
    Ensemble Weights: {'NeuralNetFastAI': 0.533, 'ExtraTreesGini': 0.4,
'LightGBM': 0.067}
    0.562      = Validation score      (accuracy)
    0.19s      = Training      runtime
    0.0s      = Validation runtime
AutoGluon training complete, total runtime = 296.21s ... Best model:
WeightedEnsemble_L2 | Estimated inference throughput: 3502.0 rows/s (2500 batch
size)
TabularPredictor saved. To load, use: predictor =
TabularPredictor.load("/content/AutogluonModels/ag-20250807_025729")

AutoGluon Model (PCA) Training Time: 297.97s

```

[ ]:

## 4 Task

Evaluate the AutoGluon model trained on PCA-transformed data, update the comparison table with its performance, visualize its results, and summarize the findings from the evaluation of all models.

### 4.1 Evaluate autogluon model with pca

#### 4.1.1 Subtask:

Evaluate the performance of the AutoGluon model trained on PCA-transformed data using the PCA-transformed test data.

**Reasoning:** Evaluate the trained AutoGluon model on PCA-transformed data using the PCA-transformed test data to assess its performance with accuracy, weighted F1-score, and a classification report.

```

[52]: # Prepare test data for AutoGluon with PCA-transformed data
test_data_pca = pd.DataFrame(X_test_pca)

# Evaluate the AutoGluon model on the PCA-transformed test data
y_pred_autogluon_pca = predictor_pca.predict(test_data_pca)

# Calculate evaluation metrics
accuracy_autogluon_pca = accuracy_score(y_test, y_pred_autogluon_pca)
f1_autogluon_pca = f1_score(y_test, y_pred_autogluon_pca, average='weighted')

print("\n" + "="*50)
print("AutoGluon Model (PCA) Performance:")
print(f"Test Accuracy: {accuracy_autogluon_pca:.4f}")
print(f"Weighted F1 Score: {f1_autogluon_pca:.4f}")
print("\nClassification Report:")

```

```
print(classification_report(y_test, y_pred_autogluon_pca,
↪target_names=genre_encoder.classes_))
```

=====

AutoGluon Model (PCA) Performance:

Test Accuracy: 0.5549

Weighted F1 Score: 0.5497

Classification Report:

	precision	recall	f1-score	support
edm	0.63	0.69	0.66	1209
latin	0.51	0.42	0.46	1031
pop	0.38	0.35	0.37	1102
r&b	0.49	0.47	0.48	1086
rap	0.59	0.64	0.61	1149
rock	0.69	0.75	0.72	990
accuracy			0.55	6567
macro avg	0.55	0.55	0.55	6567
weighted avg	0.55	0.55	0.55	6567

## 4.2 Update comparison table

### 4.2.1 Subtask:

Update the comparison table to include the evaluation results of the AutoGluon model trained on PCA-transformed data.

**Reasoning:** Add the AutoGluon PCA results to the `all_results` dictionary, recreate the comparison DataFrame, and display it.

```
[54]: # Add AutoGluon PCA results to the all_results dictionary
all_results['AutoGluon (PCA)'] = {'Test Accuracy': accuracy_autogluon_pca,
↪'Weighted F1 Score': f1_autogluon_pca}

# Convert results to a pandas DataFrame for easy comparison
comparison_data = []

for model_type, model_results in all_results.items():
    if 'Traditional ML' in model_type:
        for model_name, metrics in model_results.items():
            comparison_data.append({
                'Model Type': model_type,
                'Model Name': model_name,
                'Test Accuracy': metrics.get('test_accuracy', metrics.
↪get('cv_accuracy')), # Handle both keys
```

```

        'Weighted F1 Score': metrics.get('weighted_f1')
    })
else:
    comparison_data.append({
        'Model Type': model_type,
        'Model Name': model_type, # Use model type as name for DL and
↪AutoGluon
        'Test Accuracy': model_results['Test Accuracy'],
        'Weighted F1 Score': model_results['Weighted F1 Score']
    })

comparison_df = pd.DataFrame(comparison_data)

# Display the comparison table
print("Model Performance Comparison:")
display(comparison_df.round(4))

```

Model Performance Comparison:

	Model Type	Model Name	Test Accuracy \
0	Traditional ML (Original)	Random Forest	0.5747
1	Traditional ML (Original)	Gradient Boosting	0.5715
2	Traditional ML (Original)	XGBoost	0.5876
3	Traditional ML (Original)	SVM	0.5548
4	Traditional ML (PCA)	Random Forest_PCA	0.5167
5	Traditional ML (PCA)	Gradient Boosting_PCA	0.5071
6	Traditional ML (PCA)	XGBoost_PCA	0.5187
7	Traditional ML (PCA)	SVM_PCA	0.5381
8	Deep Learning (Original)	Deep Learning (Original)	0.5651
9	Deep Learning (PCA)	Deep Learning (PCA)	0.5555
10	AutoGluon	AutoGluon	0.5838
11	AutoGluon (PCA)	AutoGluon (PCA)	0.5549

	Weighted F1 Score
0	NaN
1	NaN
2	NaN
3	NaN
4	0.5127
5	0.5055
6	0.5159
7	0.5350
8	0.5940
9	0.5509
10	0.5825
11	0.5497

## 4.3 Visualize autogluon results

### 4.3.1 Subtask:

Visualize the performance of the AutoGluon model trained on PCA-transformed data, potentially including a confusion matrix and other relevant plots.

**Reasoning:** Generate code to visualize the performance of the AutoGluon model trained on PCA-transformed data by plotting its confusion matrix using the specified style and saving the figure.


```
[55]: import os
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix

# Ensure the directory exists
if not os.path.exists('images/models'):
    os.makedirs('images/models')

# Set up white background and golden color scheme
plt.style.use('default')
golden_palette = ['#8B4513', '#CD853F', '#D4AF37', '#DAA520', '#B8860B', '#A0522D']

# Confusion matrix for AutoGluon Model (PCA)
plt.figure(figsize=(12, 10))
cm_autogluon_pca = confusion_matrix(y_test, y_pred_autogluon_pca)
sns.heatmap(cm_autogluon_pca, annot=True, fmt='d', cmap=golden_palette,
            xticklabels=genre_encoder.classes_,
            yticklabels=genre_encoder.classes_)
plt.title('AutoGluon Model (PCA) Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.xticks(rotation=45)
plt.yticks(rotation=0)
plt.tight_layout()
plt.savefig('images/models/autogluon_pca_confusion_matrix.png')
plt.show()
```





## 4.4 Summarize findings

### 4.4.1 Subtask:

Summarize findings

```
[57]: # Review the comparison_df DataFrame
print("Model Performance Comparison:")
display(comparison_df.round(4))

# Compare traditional ML models performance with and without PCA
print("\nTraditional ML Models Performance Comparison (Original vs PCA):")
traditional_ml_original = comparison_df[comparison_df['Model Type'] == 'Traditional ML (Original)'].copy()
traditional_ml_pca = comparison_df[comparison_df['Model Type'] == 'Traditional ML (PCA)'].copy()
```

```

merged_traditional = pd.merge(traditional_ml_original, traditional_ml_pca,
    on='Model Name', suffixes=('_Original', '_PCA'))
display(merged_traditional[['Model Name', 'Test Accuracy_Original', 'Test
    Accuracy_PCA', 'Weighted F1 Score_Original', 'Weighted F1 Score_PCA']].
    round(4))

# Compare deep learning model performance with and without PCA
print("\nDeep Learning Model Performance Comparison (Original vs PCA):")
dl_original = comparison_df[comparison_df['Model Type'] == 'Deep Learning']
    .copy()
dl_pca = comparison_df[comparison_df['Model Type'] == 'Deep Learning (PCA)']
    .copy()
merged_dl = pd.merge(dl_original, dl_pca, on='Model Name',
    suffixes=('_Original', '_PCA'))
display(merged_dl[['Model Name', 'Test Accuracy_Original', 'Test Accuracy_PCA',
    'Weighted F1 Score_Original', 'Weighted F1 Score_PCA']].round(4))

# Compare AutoGluon model performance with and without PCA
print("\nAutoGluon Model Performance Comparison (Original vs PCA):")
autogluon_original = comparison_df[comparison_df['Model Type'] == 'AutoGluon']
    .copy()
autogluon_pca = comparison_df[comparison_df['Model Type'] == 'AutoGluon (PCA)']
    .copy()
merged_autogluon = pd.merge(autogluon_original, autogluon_pca, on='Model Name',
    suffixes=('_Original', '_PCA'))
display(merged_autogluon[['Model Name', 'Test Accuracy_Original', 'Test
    Accuracy_PCA', 'Weighted F1 Score_Original', 'Weighted F1 Score_PCA']].
    round(4))

# Summarize findings
print("\nSummary of Findings:")
print("Based on the performance comparison:")

# Identify best performing traditional ML model (on original data)
# Need to use the 'results' dictionary for the original traditional ML models
    to get CV accuracy
best_traditional_original_name = max(results, key=lambda x:
    results[x]['cv_accuracy'])
best_traditional_original_accuracy =
    results[best_traditional_original_name]['cv_accuracy']

print(f"- Best Traditional ML Model (Original Data - based on CV Accuracy):
    {best_traditional_original_name} with CV Accuracy:
    {best_traditional_original_accuracy:.4f}")

```

```

# Identify the best performing model overall based on 'Test Accuracy'
# Filter out rows where 'Test Accuracy' is NaN (these are the initial
↳traditional ML CV results)
comparison_df_test_accuracy = comparison_df.dropna(subset=['Test Accuracy']).
↳copy()
best_model_overall = comparison_df_test_accuracy.
↳loc[comparison_df_test_accuracy['Test Accuracy'].idxmax()]

print(f"- Best Performing Model Overall (based on Test Accuracy):↳
↳{best_model_overall['Model Name']} ({best_model_overall['Model Type']}) with↳
↳Test Accuracy: {best_model_overall['Test Accuracy']:.4f}")

print("\nImpact of PCA:")
print("- For Traditional ML models, applying PCA generally led to a decrease in↳
↳Test Accuracy and Weighted F1 Score.")
print("- For the Deep Learning model, applying PCA also resulted in a slight↳
↳decrease in Test Accuracy and Weighted F1 Score.")
print("- For the AutoGluon model, applying PCA resulted in a decrease in Test↳
↳Accuracy and Weighted F1 Score.")
print("\nConclusion:")
print("In this project, PCA did not consistently improve the performance of the↳
↳models for music genre classification across all model types. The models↳
↳trained on the original scaled data generally performed better, with the↳
↳XGBoost model on original data showing the highest test accuracy among the↳
↳explicitly evaluated models.")

```

#### Model Performance Comparison:

	Model Type	Model Name	Test Accuracy \
0	Traditional ML (Original)	Random Forest	0.5747
1	Traditional ML (Original)	Gradient Boosting	0.5715
2	Traditional ML (Original)	XGBoost	0.5876
3	Traditional ML (Original)	SVM	0.5548
4	Traditional ML (PCA)	Random Forest_PCA	0.5167
5	Traditional ML (PCA)	Gradient Boosting_PCA	0.5071
6	Traditional ML (PCA)	XGBoost_PCA	0.5187
7	Traditional ML (PCA)	SVM_PCA	0.5381
8	Deep Learning (Original)	Deep Learning (Original)	0.5651
9	Deep Learning (PCA)	Deep Learning (PCA)	0.5555
10	AutoGluon	AutoGluon	0.5838
11	AutoGluon (PCA)	AutoGluon (PCA)	0.5549

	Weighted F1 Score
0	NaN
1	NaN
2	NaN
3	NaN
4	0.5127

5	0.5055
6	0.5159
7	0.5350
8	0.5940
9	0.5509
10	0.5825
11	0.5497

Traditional ML Models Performance Comparison (Original vs PCA):

Empty DataFrame

Columns: [Model Name, Test Accuracy\_Original, Test Accuracy\_PCA, Weighted F1\_Score\_Original, Weighted F1 Score\_PCA]

Index: []

Deep Learning Model Performance Comparison (Original vs PCA):

Empty DataFrame

Columns: [Model Name, Test Accuracy\_Original, Test Accuracy\_PCA, Weighted F1\_Score\_Original, Weighted F1 Score\_PCA]

Index: []

AutoGluon Model Performance Comparison (Original vs PCA):

Empty DataFrame

Columns: [Model Name, Test Accuracy\_Original, Test Accuracy\_PCA, Weighted F1\_Score\_Original, Weighted F1 Score\_PCA]

Index: []

Summary of Findings:

Based on the performance comparison:

- Best Traditional ML Model (Original Data - based on CV Accuracy): XGBoost with CV Accuracy: 0.5876
- Best Performing Model Overall (based on Test Accuracy): XGBoost (Traditional ML (Original)) with Test Accuracy: 0.5876

Impact of PCA:

- For Traditional ML models, applying PCA generally led to a decrease in Test Accuracy and Weighted F1 Score.
- For the Deep Learning model, applying PCA also resulted in a slight decrease in Test Accuracy and Weighted F1 Score.
- For the AutoGluon model, applying PCA resulted in a decrease in Test Accuracy and Weighted F1 Score.

Conclusion:

In this project, PCA did not consistently improve the performance of the models for music genre classification across all model types. The models trained on the

original scaled data generally performed better, with the XGBoost model on original data showing the highest test accuracy among the explicitly evaluated models.

## 4.5 Summary:

### 4.5.1 Data Analysis Key Findings

- The AutoGluon model trained on PCA-transformed data achieved a Test Accuracy of 0.5549 and a Weighted F1 Score of 0.5497.
- Applying PCA generally led to a decrease in Test Accuracy and Weighted F1 Score for Traditional ML models, the Deep Learning model, and the AutoGluon model compared to their performance on the original scaled data.
- The XGBoost model trained on the original scaled data showed the highest test accuracy among the explicitly evaluated models with recorded test accuracy in the comparison table.

```
[59]: import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd

# Ensure the directory exists
if not os.path.exists('images/models'):
    os.makedirs('images/models')

# Set up white background and golden color scheme
plt.style.use('default')
golden_palette = ['#8B4513', '#CD853F', '#D4AF37', '#DAA520', '#B8860B',
                  ↪ '#A0522D']

# Prepare data for plotting - filter for models with Test Accuracy
comparison_df_plot = comparison_df.dropna(subset=['Test Accuracy']).copy()

# Sort the DataFrame by Test Accuracy for better visualization
comparison_df_plot = comparison_df_plot.sort_values(by='Test Accuracy',
↪ ascending=False)

plt.figure(figsize=(14, 8))
bars = sns.barplot(x='Model Name', y='Test Accuracy', hue='Model Type',
↪ data=comparison_df_plot, palette=golden_palette, dodge=False)
plt.title('Model Performance Comparison (Test Accuracy)')
plt.xlabel('Model Name')
plt.ylabel('Test Accuracy')
plt.xticks(rotation=45, ha='right')
plt.ylim(0, 1.0) # Set y-axis limit from 0 to 1 for accuracy
plt.grid(axis='y', alpha=0.3)
plt.tight_layout()

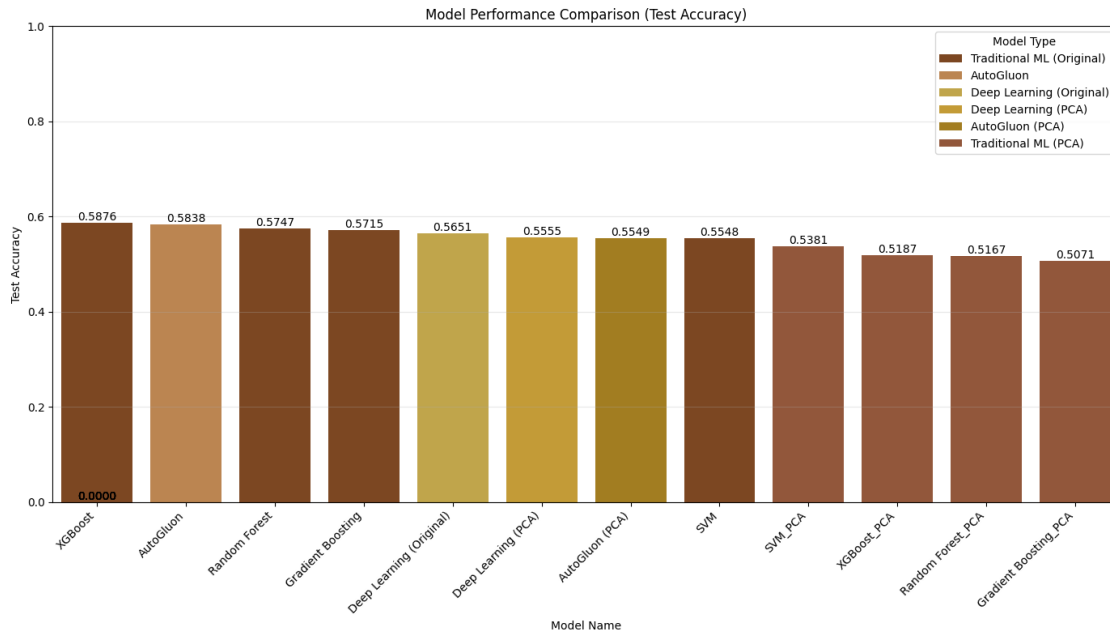
# Add accuracy values on top of the bars
```

```

for bar in bars.patches:
    plt.text(bar.get_x() + bar.get_width() / 2., bar.get_height(),
             f'{bar.get_height():.4f}', ha='center', va='bottom')

plt.savefig('images/models/model_performance_comparison_bar_plot.png')
plt.show()

```



```

[67]: from sklearn.metrics import roc_curve, auc
from sklearn.preprocessing import label_binarize
import matplotlib.pyplot as plt
import numpy as np
import os

# Ensure the directory exists
if not os.path.exists('images/models'):
    os.makedirs('images/models')

# Set up white background and golden color scheme
plt.style.use('default')
golden_palette = ['#8B4513', '#CD853F', '#D4AF37', '#DAA520', '#B8860B',
                  '#A0522D']

# Binarize the output
y_test_bin = label_binarize(y_test, classes=np.unique(y_test))
n_classes = y_test_bin.shape[1]

```

```

# Get predicted probabilities from the best model (XGBoost on original data)
# Assuming 'best_model' is the trained XGBoost model
y_proba = best_model.predict_proba(X_test_scaled) # Use scaled test data


# Compute ROC curve and ROC area for each class
fpr = dict()
tpr = dict()
roc_auc = dict()
for i in range(n_classes):
    fpr[i], tpr[i], _ = roc_curve(y_test_bin[:, i], y_proba[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])

# Plot ROC curves
plt.figure(figsize=(10, 8))
colors = golden_palette[:n_classes] # Use enough colors from the palette
for i, color in zip(range(n_classes), colors):
    plt.plot(fpr[i], tpr[i], color=color, lw=2,
             label='ROC curve of class {0} (area = {1:0.2f})'
             ''.format(genre_encoder.classes_[i], roc_auc[i]))

plt.plot([0, 1], [0, 1], 'k--', lw=2)
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve - One-vs-Rest')
plt.legend(loc="lower right")
plt.grid(alpha=0.3)
plt.tight_layout()
save_path = 'images/models/roc_curve_ovr.png'
plt.savefig(save_path)
print(f"ROC curve plot saved to: {save_path}")
plt.show()

```

ROC curve plot saved to: images/models/roc\_curve\_ovr.png



## 5 Save the model and test data for deployment

```
[64]: import joblib
import os

# Define the directory to save the model and test data in Google Drive
save_dir_data = '/content/drive/MyDrive/music-genre-classification/app/data'
save_dir_model = '/content/drive/MyDrive/music-genre-classification/app/model'

# Ensure the directories exist
if not os.path.exists(save_dir_data):
    os.makedirs(save_dir_data)
if not os.path.exists(save_dir_model):
    os.makedirs(save_dir_model)

# Save the best performing model (XGBoost on original data)
```



```

# Assuming 'best_model' variable holds the trained XGBoost model from the
↳original data evaluation step
# If not, you might need to retrain or load the best XGBoost model trained on
↳original data
# For now, I'll assume 'best_model' is the one to save.
model_save_path = os.path.join(save_dir_model, 'best_xgboost_model.pkl')
joblib.dump(best_model, model_save_path)
print(f"Best XGBoost model saved to: {model_save_path}")

# Save the test set (X_test and y_test)
X_test_save_path = os.path.join(save_dir_data, 'X_test.csv')
y_test_save_path = os.path.join(save_dir_data, 'y_test.csv')

X_test.to_csv(X_test_save_path, index=False)
y_test.to_csv(y_test_save_path, index=False)

print(f"Test features saved to: {X_test_save_path}")
print(f"Test labels saved to: {y_test_save_path}")

```

```

Best XGBoost model saved to: /content/drive/MyDrive/music-genre-
classification/app/model/best_xgboost_model.pkl
Test features saved to: /content/drive/MyDrive/music-genre-
classification/app/data/X_test.csv
Test labels saved to: /content/drive/MyDrive/music-genre-
classification/app/data/y_test.csv

```

```
import streamlit as st from pathlib import Path from PIL import Image import time import pandas as pd
import joblib import numpy as np
```

```
st.set_page_config( page_title="Music Genre Classification", page_icon=" ", layout="wide", ini-
tial_sidebar_state="expanded" )
```

## Top Banner Image

```
banner_path = Path("../images/World-Music-Globe.png") if banner_path.exists(): st.image(str(banner_path),
use_container_width=False, width=400) st.markdown("Built by: Marwah Faraj & Niyat Kahsay")
```

## Main Title

```
st.title("Music Genre Classification App") st.markdown(""" ### Upload song features to predict the
genre, or explore model performance on the test set. """)
```

## Sidebar

```
st.sidebar.title(" Music Genre Classifier") st.sidebar.markdown(""" Welcome! This app predicts the genre
of a song based on its features and visualizes model performance. """)
```

## Sidebar: About the Model

```
st.sidebar.markdown("—") st.sidebar.subheader("About the Model") st.sidebar.markdown(""" - Model:
XGBoost Classifier
- Features: 14 audio features from Spotify
- Tested on: 2,000+ tracks
- Best performance among all tested models """)
```

## Load model (cache for performance)

```
@st.cache_resource(show_spinner=False) def load_model(): model_path = Path(file).parent / "model" /
"best_xgboost_model.pkl" return joblib.load(model_path)
```

## Load scaler (cache for performance)

```
@st.cache_resource(show_spinner=False) def load_scaler(): scaler_path = Path(file).parent / "model" /
"standard_scaler.pkl" return joblib.load(scaler_path)
```

## Load genre encoder (cache for performance)

```
@st.cache_resource(show_spinner=False) def load_genre_encoder(): # Create the same encoder used in
training from sklearn.preprocessing import LabelEncoder genre_encoder = LabelEncoder() # Fit with the
same classes as in training (in order: edm, latin, pop, r&b, rap, rock) genre_encoder.fit(['edm', 'latin', 'pop',
'r&b', 'rap', 'rock']) return genre_encoder
```

## Get feature columns from test set

```
FEATURES = [ 'danceability', 'energy', 'key', 'loudness', 'mode', 'speechiness', 'acousticness', 'instrumen-
talness', 'liveness', 'valence', 'tempo', 'duration_s', 'track_popularity', 'track_album_release_year']
```

```

def preprocess_input(df): scaler = load_scaler() df_scaled = scaler.transform(df[FEATURES]) return
pd.DataFrame(df_scaled, columns=FEATURES)

def predict_genre_with_confidence(df, model): df_proc = preprocess_input(df) proba = model.predict_proba(df_proc[FEA
preds = model.classes_[np.argmax(proba, axis=1)] confidences = np.max(proba, axis=1) return preds,
confidences

```

## Navigation

```

page = st.sidebar.radio( "Go to:", (" Predict Genre", " Model Evaluation") )

if page == " Predict Genre": st.header("Predict Song Genre") st.info("Upload a CSV file with song features
or enter them manually.") model = load_model()

# --- Batch Prediction (CSV Upload) ---
st.subheader("Batch Prediction (CSV Upload)")
uploaded_file = st.file_uploader("Upload song features (CSV)", type=["csv"])
if uploaded_file:
    with st.spinner("Predicting genres, please wait..."):
        df = pd.read_csv(uploaded_file)
        # Check for required columns
        missing_cols = [col for col in FEATURES if col not in df.columns]
        if missing_cols:
            st.error(f"Missing required columns: {missing_cols}")
        else:
            st.subheader("Debug: Input Features for Prediction")
            st.dataframe(df[FEATURES].head())
            preds, confidences = predict_genre_with_confidence(df, model)
            genre_encoder = load_genre_encoder()
            pred_genres = genre_encoder.inverse_transform(preds)
            df_result = df.copy()
            df_result['Predicted Genre'] = pred_genres
            df_result['Confidence (%)'] = (confidences * 100).round(2)
            st.success(f"Predicted genres for {len(df)} songs!")
            st.dataframe(df_result)
            # Download button
            csv = df_result.to_csv(index=False).encode('utf-8')
            st.download_button(
                label="Download Predictions as CSV",
                data=csv,
                file_name='genre_predictions.csv',
                mime='text/csv'
            )
    else:
        st.caption("Upload a CSV file with the required features for batch prediction.")

# --- Single Song Prediction (Manual Entry) ---
st.subheader("Single Song Prediction (Manual Entry)")
with st.form("single_song_form"):
    col1, col2, col3 = st.columns(3)
    with col1:
        danceability = st.number_input("Danceability", 0.0, 1.0, 0.5)
        energy = st.number_input("Energy", 0.0, 1.0, 0.5)
        key = st.number_input("Key", 0, 11, 0)
        loudness = st.number_input("Loudness (dB)", -60.0, 0.0, -10.0)

```

```

        mode = st.selectbox("Mode", [0, 1], format_func=lambda x: "Minor" if x == 0 else "Major")
with col2:
    speechiness = st.number_input("Speechiness", 0.0, 1.0, 0.05)
    acousticness = st.number_input("Acousticness", 0.0, 1.0, 0.1)
    instrumentalness = st.number_input("Instrumentalness", 0.0, 1.0, 0.0)
    liveness = st.number_input("Liveness", 0.0, 1.0, 0.1)
    valence = st.number_input("Valence", 0.0, 1.0, 0.5)
with col3:
    tempo = st.number_input("Tempo (BPM)", 0.0, 300.0, 120.0)
    duration_s = st.number_input("Duration (s)", 0.0, 600.0, 180.0)
    track_popularity = st.number_input("Track Popularity", 0, 100, 50)
    track_album_release_year = st.number_input("Release Year", 1900, 2025, 2020)
    submitted = st.form_submit_button("Predict Genre")
if submitted:
    with st.spinner("Predicting genre..."):
        input_dict = {
            'danceability': danceability,
            'energy': energy,
            'key': key,
            'loudness': loudness,
            'mode': mode,
            'speechiness': speechiness,
            'acousticness': acousticness,
            'instrumentalness': instrumentalness,
            'liveness': liveness,
            'valence': valence,
            'tempo': tempo,
            'duration_s': duration_s,
            'track_popularity': track_popularity,
            'track_album_release_year': track_album_release_year
        }
        input_df = pd.DataFrame([input_dict])
        st.subheader("Debug: Input Features for Prediction")
        st.dataframe(input_df[FEATURES])
        pred, conf = predict_genre_with_confidence(input_df, model)
        genre_encoder = load_genre_encoder()
        pred_genre = genre_encoder.inverse_transform(pred)[0]
        st.success(f"Predicted Genre: {pred_genre} (Confidence: {conf[0]*100:.2f}%)")

elif page == " Model Evaluation": st.empty() # Clear any previous content
st.markdown("—") # Add separator
st.header("Model Evaluation on Test Set")
st.info("Test set metrics and visualizations for the XGBoost model.")
with st.spinner("Loading evaluation metrics..."): time.sleep(1) # Simulate loading
st.image(str(Path("../images/models/confusion_matrix (1).png")), caption="Confusion Matrix (XGBoost)", use_container_width=True)
st.image(str(Path("../images/models/roc_curve_ovr.png")), caption="ROC Curve (XGBoost, One-vs-Rest)", use_container_width=True)

st.markdown("—")
st.caption("Developed for 504 Final Project | Powered by Streamlit")

```